

# Efficient Algorithms for Geometric Partial Matching

**Pankaj K. Agarwal**

Duke University, USA

[pankaj@cs.duke.edu](mailto:pankaj@cs.duke.edu)

**Hsien-Chih Chang**

Duke University, USA

[hsienchih.chang@duke.edu](mailto:hsienchih.chang@duke.edu)

**Allen Xiao**

Duke University, USA

[axiao@cs.duke.edu](mailto:axiao@cs.duke.edu)

## 1 Abstract

Let  $A$  and  $B$  be two point sets in the plane of sizes  $r$  and  $n$  respectively (assume  $r \leq n$ ), and let  $k$  be a parameter. A matching between  $A$  and  $B$  is a family  $M \subseteq A \times B$  of pairs so that any point of  $A \cup B$  appears in at most one pair. Given two integers  $p, q \geq 1$ , we define the cost of  $M$  to be  $\text{cost}(M) = \sum_{(a,b) \in M} \|a - b\|_p^q$  where  $\|\cdot\|_p$  is the  $L_p$ -norm. The geometric partial matching problem asks to find the minimum-cost size- $k$  matching between  $A$  and  $B$ .

We present efficient algorithms for geometric partial matching that work for any powers of  $L_p$ -norm matching objective: An exact algorithm that runs in  $O((n + k^2) \text{polylog } n)$  time, and a  $(1 + \varepsilon)$ -approximation algorithm that runs in  $O((n + k\sqrt{k}) \text{polylog } n \cdot \log \varepsilon^{-1})$  time. Both algorithms are based on the primal-dual flow augmentation scheme; the main improvements are obtained by using dynamic data structures to achieve efficient flow augmentations. Using similar techniques, we give an exact algorithm for the planar transportation problem that runs in  $O((r^2\sqrt{n} + rn^{3/2}) \text{polylog } n)$  time. For  $r = o(\sqrt{n})$ , this algorithm is faster than the state-of-art near-quadratic time algorithm by Agarwal *et al.* [SOCG 2017].

**2012 ACM Subject Classification** Theory of computation  $\rightarrow$  Design and analysis of algorithms

**Keywords and phrases** partial matching, transportation, optimal transport, minimum-cost flow, bichromatic closest pair

**Funding** Work on this paper was supported by NSF under grants CCF-15-13816, CCF-15-46392, and IIS-14-08846, by an ARO grant W911NF-15-1-0408, and by BSF Grant 2012/229 from the U.S.-Israel Binational Science Foundation.

**Acknowledgements** We thank Haim Kaplan for discussion and suggestion to use Goldberg *et al.* [13] algorithm. We thank Debmalya Panigrahi for helpful discussions.

Lines 517

## 1 Introduction

Given two point sets  $A$  and  $B$  in  $\mathbb{R}^2$ , we consider the problem of finding the minimum-cost partial matching between  $A$  and  $B$ . Formally, suppose  $A$  has size  $r$  and  $B$  has size  $n$  where  $r \leq n$ . Let  $G(A, B)$  be the undirected complete bipartite graph between  $A$  and  $B$ , and let the cost of edge  $(a, b)$  be  $c(a, b) = \|a - b\|_p$ , for some  $1 \leq p < \infty$ . A *matching*  $M$  in  $G(A, B)$



© Pankaj K. Agarwal, Hsien-Chih Chang, Allen Xiao;  
licensed under Creative Commons License CC-BY

The 35th International Symposium on Computational Geometry (SOCG 2019).



Leibniz International Proceedings in Informatics

LIPICs Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany



is a set of edges sharing no endpoints. The *size* of  $M$  is the number of edges in  $M$ . Given  $q \geq 1$ , the cost of  $M$  is defined to be

$$\text{cost}(M) := \sum_{(a,b) \in M} \|a - b\|_p^q.$$

For a parameter  $k \leq r$ , the problem of finding the minimum-cost size- $k$  matching in  $G(A, B)$  is called the *geometric partial matching problem*. We call the corresponding problem in general bipartite graphs (with arbitrary edge costs) the *partial matching problem*.<sup>1</sup>

We also consider the following generalization of bipartite matching. Let  $\lambda : A \cup B \rightarrow \mathbb{Z}$  be a *supply-demand function* with positive value on points of  $A$  and negative value on points of  $B$ , satisfying  $\sum_{a \in A} \lambda(a) = -\sum_{b \in B} \lambda(b)$ . A *transportation map* is a function  $\tau : A \times B \rightarrow \mathbb{R}_{\geq 0}$  such that  $\sum_{b \in B} \tau(a, b) = \lambda(a)$  for all  $a \in A$  and  $\sum_{a \in A} \tau(a, b) = -\lambda(b)$  for all  $b \in B$ . We define the cost of  $\tau$  to be

$$\text{cost}(\tau) := \sum_{(a,b) \in A \times B} \|a - b\|_p^q \cdot \tau(a, b).$$

The *transportation problem* asks to compute a transportation map of minimum cost.

**Related work** Maximum-size bipartite matching is a classical problem in graph algorithms. Upper bounds include the  $O(m\sqrt{n})$  time algorithm by Hopcroft and Karp [15] and the  $O(m \min\{\sqrt{m}, n^{2/3}\})$  time algorithm by Even and Tarjan [11], where  $n$  is the number of vertices and  $m$  is the number of edges. The first improvement in over thirty years was made by Mądry [19], which uses an interior-point algorithm and runs in  $O(m^{10/7} \text{polylog } n)$  time.

The Hungarian algorithm [17] computes a minimum-cost maximum matching in a bipartite graph in roughly  $O(mn)$  time. Faster algorithms have been developed, such as the  $O(m\sqrt{n} \log(nC))$  time algorithms by Gabow and Tarjan [12] and the improved  $O(m\sqrt{n} \log C)$  time algorithm by Duan *et al.* [9] assuming that the edge costs are integral; here  $C$  is the maximum cost of an edge. Ramshaw and Tarjan [21] showed that the Hungarian algorithm can be extended to compute a minimum-cost partial matching of size  $k$  in time  $O(km + k^2 \log r)$  time. They also proposed a cost-scaling algorithm for partial matching that runs in time  $O(m\sqrt{n} \log(nC))$ , again assuming that costs are integral. By reduction to unit-capacity min-cost flow, Goldberg *et al.* [13] developed a cost-scaling algorithm for partial matching with running time  $O(m\sqrt{k} \log(kC))$ , again only for integral edge costs.

In geometric settings, the Hungarian algorithm can be implemented to compute an optimal perfect matching between  $A$  and  $B$  (assuming equal size) in time  $O(n^2 \text{polylog } n)$  [16] (see also [1, 26]). This algorithm computes an optimal size- $k$  matching in time  $O(kn \text{polylog } n)$ . Faster approximation algorithms have been developed for computing a perfect matching in geometric settings [4, 23, 26, 27]. For  $q = 1$ , the best algorithm to date by Sharathkumar and Agarwal [24] computes a  $(1 + \varepsilon)$ -approximation to the optimal perfect matching in  $O(n \text{polylog } n \cdot \varepsilon^{-O(1)})$  expected time with high probability. Their algorithm can also compute a  $(1 + \varepsilon)$ -approximate partial matching within the same time bound. For  $q > 1$ , the best known approximation algorithm to compute a perfect matching runs in  $O(n^{3/2} \text{polylog } n \log(n/\varepsilon))$  time [23]; it is not obvious how to extend this algorithm to the partial matching setting.

There is also some work on computing an optimal or near-optimal partial matching when  $B$  is fixed but  $A$  is allowed to translate and/or rotate [3, 5, 8, 22]. Here, the goal is to (i)

<sup>1</sup> Partial matching is also called *imperfect matching* or *imperfect assignment* [13, 21].

compute a (near-)optimal matching over all possible transformations of  $A$ , or (ii) to compute a set  $\mathcal{M}$  of matchings, such that for any translation/rotation  $\mathcal{T}$  of  $A$ , a (near-)optimal matching of  $\mathcal{T}(A)$  and  $B$  in  $\mathcal{M}$ .

The transportation problem can also be formulated as a minimum-cost flow problem in a graph. The strongly polynomial uncappeditated min-cost flow algorithm by Orlin [20] solves the transportation problem in  $O((m + n \log n)n \log n)$  time. Lee and Sidford [18] give a weakly polynomial algorithm that runs in  $O(m\sqrt{n} \text{polylog}(n, U))$  time, where  $U$  is the maximum amount of vertex supply-demand. Agarwal *et al.* [2] showed that Orlin's algorithm can be implemented to solve 2D transportation in time  $O(n^2 \text{polylog } n)$ . It is not known whether the 2D transportation problem can be solved in  $O(rn \text{polylog } n)$  time. By adapting the Lee-Sidford algorithm, they developed a  $(1 + \varepsilon)$ -approximation algorithm that runs in  $O(n^{3/2}\varepsilon^{-2} \text{polylog}(n, U))$  time. They also gave a Monte-Carlo algorithm that computes an  $O(\log^2(1/\varepsilon))$ -approximate solution in  $O(n^{1+\varepsilon})$  time with high probability.

**Our results** There are three main results in this paper. First in Section 2 we present an efficient algorithm for computing an optimal partial matching in  $\mathbb{R}^2$ .

► **Theorem 1.1.** *Given two point sets  $A$  and  $B$  in  $\mathbb{R}^2$  each of size at most  $n$ , a minimum-cost matching of size  $k$  between  $A$  and  $B$  can be computed in  $O((n + k^2) \text{polylog } n)$  time.*

We use *bichromatic closest pair (BCP)* data structures to implement the Hungarian algorithm efficiently, similar to Agarwal *et al.* and Kaplan *et al.* [1, 16]. But unlike their algorithms which take  $\Omega(n)$  time to find an augmenting path, we show that after  $O(n \text{polylog } n)$  preprocessing, an augmenting path can be found in  $O(k \text{polylog } n)$  time. The key is to recycle (rather than rebuild) our data structures from one augmentation to the next. We refer to this idea as the *rewinding mechanism*.

Next in Sections 3 and 4, we obtain a  $(1 + \varepsilon)$ -approximation algorithm for the geometric partial matching problem in  $\mathbb{R}^2$  by providing an efficient implementation of the unit-capacity min-cost flow algorithm by Goldberg *et al.* [13].

► **Theorem 1.2.** *Given two point sets  $A$  and  $B$  in  $\mathbb{R}^2$  each of size at most  $n$ , a  $(1 + \varepsilon)$ -approximate min-cost matching of size  $k$  between  $A$  and  $B$  can be computed in  $O((n + k\sqrt{k}) \text{polylog } n \cdot \log \varepsilon^{-1})$  time.*

The main challenge here is the set of *null vertices* which do not play any role in the augmentations, but still contribute to the size of the graph. Instead, we run the unit-capacity min-cost flow algorithm on a *shortcut graph*, circumventing all null vertices. The shortcut graph itself may have  $\Omega(n^2)$  edges, but we can represent it implicitly and use a data structure to explore this graph. As such, we can implement each augmentation in the Goldberg *et al.* algorithm in time proportional to the size of the *flow support*, which turns out to be of size  $O(k)$ .

Finally in Section 5 we present a faster algorithm for the transportation problem in  $\mathbb{R}^2$  when the two point sets are unbalanced.

► **Theorem 1.3.** *Given two point sets  $A$  and  $B$  in  $\mathbb{R}^2$  of sizes  $r$  and  $n$  respectively along with supply-demand function  $\lambda : A \cup B \rightarrow \mathbb{Z}$ , an optimal transportation map between  $A$  and  $B$  can be computed in  $O((r^2\sqrt{n} + rn^{3/2}) \text{polylog } n)$  time.*

Our result improves over the  $O(n^2 \text{polylog } n)$  time algorithm in [2] when  $r = o(\sqrt{n})$ . The algorithm uses the strongly polynomial uncappeditated minimum-cost flow algorithm

by Orlin [20], adapted for geometric costs as in Agarwal *et al.* [2]. Unlike in the case of matchings, the flow support for the transportation problem may have size  $\Omega(n)$  even when  $r$  is a constant; so naively we can no longer charge the execution time to flow support size. However, we show that most of the support arcs are of degree one and thus can be partitioned into *stars* centered at vertices of  $A$ . We describe a data structure that processes these stars in amortized  $O((r^2/\sqrt{n} + r\sqrt{n}) \text{polylog } n)$  time per augmentation.

## 2 Minimum-Cost Partial Matchings using Hungarian Algorithm

In this section, we solve the geometric partial matching problem and prove Theorem 1.1 by implementing the Hungarian algorithm for partial matching in  $O((n + k^2) \text{polylog } n)$  time.

The dual to the standard linear program for partial matching has dual variables for each vertex, called *potentials*  $\pi$ . Given potentials  $\pi$ , we can define the *reduced cost* on the edges to be  $c_\pi(v, w) := c(v, w) - \pi(v) + \pi(w)$ . Potentials  $\pi$  are *feasible* if the reduced costs are nonnegative for all edges in  $G$ . We say that an edge  $(v, w)$  is *admissible* under potentials  $\pi$  if  $c_\pi(v, w) = 0$ .

A vertex  $v$  is *matched* by matching  $M \subseteq E$  if  $v$  is the endpoint of some edge in  $M$ ; otherwise  $v$  is *unmatched*. Given a matching  $M$ , an *augmenting path*  $\Pi = (a_1, b_1, \dots, a_\ell, b_\ell)$  is an odd-length path with unmatched endpoints  $(a_1$  and  $b_\ell)$  that alternates between edges outside and inside of  $M$ . The symmetric difference  $M \oplus \Pi$  creates a new matching of size  $|M| + 1$ , called the *augmentation* of  $M$  by  $\Pi$ .

**Fast implementation of Hungarian search** The Hungarian algorithm is initialized with  $M = \emptyset$  and  $\pi = 0$ . Each iteration of the Hungarian algorithm augments  $M$  with an admissible augmenting path  $\Pi$ , discovered using a procedure called the *Hungarian search*. The algorithm terminates after  $k$  augmentations, when  $|M| = k$ ; Ramshaw and Tarjan [21] showed that  $M$  is guaranteed to be an optimal partial matching.

The Hungarian search grows a set of *reachable vertices*  $S$  from all unmatched  $v \in A$  using augmenting paths of admissible edges. Initially,  $S$  is the set of unmatched vertices in  $A$ . Let the *frontier* of  $S$  be the edges in  $(A \cap S) \times (B \setminus S)$ .  $S$  is grown by *relaxing* an edge  $(a, b)$  in the frontier: adding  $b$  into  $S$ , and also modifying potentials to make  $(a, b)$  admissible, preserve  $c_\pi$  on other edges within  $S$ , and keep  $\pi$  feasible on edges outside of  $S$ . Specifically, the algorithm relaxes the minimum-reduced-cost frontier edge  $(a, b)$ , and then raises  $\pi(v)$  by  $c_\pi(a, b)$  for all  $v \in S$ . If  $b$  is already matched, then we also relax the matching edge  $(a', b)$  and add  $a'$  into  $S$ . The search finishes when  $b$  is unmatched, and an admissible augmenting path now can be recovered.

### 2.1 Fast implementation of Hungarian search

In the geometric setting, we find the min-reduced-cost frontier edge using a dynamic *bichromatic closest pair* (BCP) data structure, as observed in [2, 26]. Given two point sets  $P$  and  $Q$  in the plane and a weight function  $\omega : P \cup Q \rightarrow \mathbb{R}$ , the BCP is two points  $a \in P$  and  $b \in Q$  minimizing the additively weighted distance  $c(a, b) - \omega(a) + \omega(b)$ . Thus, a minimum reduced-cost frontier edge is precisely the BCP of point sets  $P = A \cap S$  and  $Q = B \setminus S$ , with  $\omega = \pi$ . Note that the “state” of this BCP is parameterized by  $S$  and  $\pi$ .

The dynamic BCP data structure by Kaplan *et al.* [16] supports point insertions and deletions in  $O(\text{polylog } n)$  time and answers queries in  $O(\log^2 n)$  time for our setting. Outside of potential updates, each relaxation in the Hungarian search requires one query, one deletion, and at most one insertion. As  $|M| \leq k$  throughout, there are at most  $2k$  relaxations in

each Hungarian search, and the BCP can be used to implement each Hungarian search in  $O(k \text{ polylog } n)$  time. Explained shortly, there is an existing technique to handle potential updates without performing BCP updates for each one.

**Rewinding mechanism.** We cannot afford to take  $O(n \text{ polylog } n)$  time to initialize the BCP data structure at the beginning of every Hungarian search beyond the first one. To resolve the issue, observe that exactly one vertex of  $A$  is newly matched after an augmentation. Thus (modulo potential changes), given the initial state of the BCP at the  $i$ -th Hungarian search, we can obtain the initial state for the  $(i+1)$ -th with a single BCP deletion operation. Suppose we log the sequence of points added to  $S$  in the  $i$ -th Hungarian search. Then, at the start of the  $(i+1)$ -th Hungarian search, we can *rewind* this log by undoing each BCP update operation in reverse order of the log to obtain the initial state of the  $i$ -th BCP. With one additional BCP delete, we have the initial state for the  $(i+1)$ -th BCP. The number of points in the log is  $O(k)$ , bounded by the number of relaxations per Hungarian search. Thus, in  $O(k \text{ polylog } n)$  time we can recover the initial BCP data structure for each Hungarian search beyond the first. We refer to this procedure as the *rewinding mechanism*.

As for potential updates, we adapt an update batching trick from Vaidya [26] to work under the rewinding mechanism, so that the time spent on potential updates per Hungarian search is  $O(k)$ . See the full version for details. Putting everything together, our implementation of the Hungarian algorithm runs in  $O((n + k^2) \text{ polylog } n)$  time. This proves Theorem 1.1.

### 3 Approximating Min-Cost Partial Matching through Cost-Scaling

In this section we present a  $(1 + \varepsilon)$ -approximation algorithm for computing a geometric partial matching that runs in time  $O((n + k\sqrt{k}) \text{ polylog } n \log(1/\varepsilon))$ .

After introducing the necessary terminologies in Section 3.1, we reduce the partial matching problem to computing an approximate minimum-cost flow on a unit-capacity reduction network in Section 3.2. In Section 3.3 we prove a high-level overview of the cost-scaling algorithm, executed on the reduction network. We postpone the fast implementation using dynamic data structures to Section 4.

#### 3.1 Preliminaries on Network Flows

Due to the space restriction, we omit the definitions of standard network flow theory terminologies from the main text. For a reference see Appendix B.1, or any texts on network flows [13, 20]. We emphasize that a directed graph  $G = (V, E)$  is augmented by edge costs  $c$  and capacities  $u$ , and a supply-demand function  $\phi : V \rightarrow \mathbb{Z}$ . We call the positive values of  $\phi(v)$  *supply*, and the negative values of  $\phi(v)$  *demand*. We assume that the sum of supplies and demands is equal to zero. For Sections 3 and 4, we assume the capacity of every edge in  $G$  is 1, i.e.  $G$  is a unit capacity graph. A *network*  $N = (V, \vec{E})$  turns each edge in  $E$  into a pair of *arcs*  $v \rightarrow w$  and  $w \rightarrow v$  in arc set  $\vec{E}$ . With the unit-capacity assumption on the network, all the pseudoflows in this section take integer values. The *support* of a pseudoflow  $f$  in  $N$ , is the set of arcs with positive flows:  $\text{supp}(f) := \{v \rightarrow w \in \vec{E} \mid f(v \rightarrow w) > 0\}$ . Given a pseudoflow  $f$ , we define the *imbalance* of a vertex to be  $\phi_f(v) := \phi(v) + \sum_{w \rightarrow v \in \vec{E}} f(w \rightarrow v) - \sum_{v \rightarrow w \in \vec{E}} f(v \rightarrow w)$ . If all vertices are *balanced*, the pseudoflow is a *circulation*. The *cost* of a pseudoflow is defined to be

$$\text{cost}(f) := \sum_{v \rightarrow w \in \text{supp}(f)} c(v \rightarrow w) \cdot f(v \rightarrow w).$$

190 The *minimum-cost flow problem (MCF)* asks to find a circulation of minimum cost inside a  
 191 given directed graph.

192 «**Move to when used**» The *bottleneck value* of a flow is

$$193 \quad \text{bottleneck}(f) := \max_{v \rightarrow w \in \text{supp}(f)} c(v \rightarrow w).$$

194 The *minimum bottleneck value* for a network is the minimum bottleneck value over all  
 195 circulations.

196 **LP-duality and admissibility.** We solve the min-cost flow problem using primal-dual al-  
 197 gorithms. Let  $G = (V, E)$  be a given directed graph with the corresponding network  
 198  $N = (V, \vec{E}, c, u, \phi)$ . Formally, the *potentials*  $\pi(v)$  are the variables of the linear program dual  
 199 to the standard linear program for the min-cost flow problem with variables  $f(v \rightarrow w)$  for each  
 200 directed edge in  $E$ . Assignments to the primal variables satisfying the capacity constraints  
 201 extend naturally into a pseudoflow on the network  $N$ . Let  $G_f = (V, \vec{E}_f)$  be the residual  
 202 graph under pseudoflow  $f$ . The *reduced cost* of an arc  $v \rightarrow w$  in  $\vec{E}_f$  with respect to  $\pi$  is defined  
 203 as

$$204 \quad c_\pi(v \rightarrow w) := c(v \rightarrow w) - \pi(v) + \pi(w).$$

205 Notice that the cost function  $c_\pi$  is also antisymmetric.

206 The *dual feasibility constraint* says that  $c_\pi(v \rightarrow w) \geq 0$  holds for every directed edge  $(v, w)$   
 207 in  $E$ ; potentials  $\pi$  which satisfy this constraint are said to be *feasible*. Suppose we relax the  
 208 dual feasibility constraint to allow some small violation in the value of  $c_\pi(v \rightarrow w)$ . We say that  
 209 a pair of pseudoflow  $f$  and potential  $\pi$  is  *$\theta$ -optimal* [6, 25] if  $c_\pi(v \rightarrow w) \geq -\theta$  for every residual  
 210 arc  $v \rightarrow w$  in  $\vec{E}_f$ . Pseudoflow  $f$  is  *$\theta$ -optimal* if it is  $\theta$ -optimal with respect to some potentials  
 211  $\pi$ ; potential  $\pi$  is  *$\theta$ -optimal* if it is  $\theta$ -optimal with respect to some pseudoflow  $f$ . Given a  
 212 pseudoflow  $f$  and potentials  $\pi$ , a residual arc  $v \rightarrow w$  in  $\vec{E}_f$  is *admissible* if  $c_\pi(v \rightarrow w) \leq 0$ . We  
 213 say that a pseudoflow  $g$  in  $G_f$  is *admissible* if all support arcs of  $g$  on  $G_f$  are admissible;  
 214 in other words,  $g(v \rightarrow w) > 0$  holds only on admissible arcs  $v \rightarrow w$ . Throughout the rest of  
 215 the paper we make use of the property that performing an admissible flow augmentation  
 216 preserves  $\theta$ -optimality. To our knowledge this is folklore; see Lemma B.1 for a proof.

## 217 3.2 Reduction to Unit-Capacity Min-Cost Flow Problem

218 Here, we reduce the min-cost partial matching problem to finding a  $\theta$ -optimal circulation in  
 219 a unit-capacity min-cost flow instance.

220 **Additive approximation.** Given a bipartite graph  $G = (A, B, E_0)$  for the geometric partial  
 221 matching problem with cost function  $c$ , we construct the *reduction graph*  $H$  as follows: Direct  
 222 the edges in  $E_0$  from  $A$  to  $B$ , and assign each directed edge with capacity 1. Now add a  
 223 dummy vertex  $s$  with directed edges to all vertices in  $A$ , and add a dummy vertex  $t$  with  
 224 directed edges from all vertices in  $B$ ; each edge added has cost 0 and capacity 1. Assign  
 225 vertex  $s$  with supply  $k$  and vertex  $t$  with demand  $-k$ ; the rest of the vertices in  $H$  have zero  
 226 supply-demand. We call the network corresponding to  $H$  the *reduction network*  $N_H$ . It is  
 227 straightforward to show that any integer circulation  $f$  on  $N_H$  uses exactly  $k$  of the  $A$ -to- $B$   
 228 arcs, which correspond to the edges of a size- $k$  matching  $M_f$ . Notice that the cost of the  
 229 circulation  $f$  is equal to the cost of the corresponding matching  $M_f$ .

230 First we show that the number of arcs used by any integer pseudoflow in  $N_H$  is bounded  
 231 by the excess of the pseudoflow.



232 ► **Lemma 3.1.** *The size of  $\text{supp}(f)$  is at most  $3k$  for any integer circulation  $f$  in reduction*  
 233 *network  $N_H$ . As a corollary, the number of residual backward arcs is at most  $3k$ .*

234 Using the bound on the support size, we show that a  $\theta$ -optimal integral circulation gives  
 235 an additive  $O(k\theta)$ -approximation to the MCF problem.

236 ► **Lemma 3.2.** *Let  $f$  be a  $\theta$ -optimal integer circulation in  $N_H$ , and  $f^*$  be an optimal integer*  
 237 *circulation for  $N_H$ . Then,  $\text{cost}(f) \leq \text{cost}(f^*) + 6k\theta$ .*

238 **Multiplicative approximation.** We employ a technique from Sharathkumar and Agarwal [23]  
 239 to convert the additive approximation into a multiplicative one. The reduction does not  
 240 work out of the box, as they were tackling a similar but different problem on geometric  
 241 transportations. See Appendix B.2 for details.

242 ► **Lemma 3.3.** *Computing  $(1 + \varepsilon)$ -approximate geometric partial matching reduces to the*  
 243 *following problem in  $O(n \text{ polylog } n)$  time: Given a fixed  $\alpha$ , compute an  $(\alpha\varepsilon/6k)$ -optimal*  
 244 *circulation on a reduction network  $N$  with min-bottleneck value at most  $n^q\alpha$ .*

### 245 3.3 High-Level Description of Cost-Scaling Algorithm

249 The main body of our algorithm is an implementation of the unit-capacity min-cost flow  
 250 algorithm of Goldberg *et al.* [13]. Their algorithm is based on the *cost-scaling* technique,  
 251 originally due to Goldberg and Tarjan [14]. The algorithm finds  $\theta$ -optimal circulations for  
 252 geometrically shrinking values of  $\theta$ . Each fixed value of  $\theta$  is called a *cost scale*. Once  $\theta$  is  
 253 sufficiently small, the  $\theta$ -optimal flow is a suitable approximation according to Lemma 3.3.<sup>2</sup>

254 In what follows, we use  $\theta^*$  to denote the “target” scale of the cost-scaling (that is, we seek  
 255 a  $\theta^*$ -optimal circulation). The cost-scaling algorithm initializes the flow  $f$  and the potential  
 256  $\pi$  to be zero.

257 **«Unclear; rewrite»** The initial value of the scale parameter, denoted  $\theta_0$ , is set to  $C$  by  
 258 the original algorithm [13], but this does not work well with the multiplicative approximation  
 259 reduction (Lemma 3.3). Instead, we use

$$260 \quad \theta_0 \in 2 \cdot \{x \in \mathbb{R}_{\geq 0} \mid \text{there exists a circulation supported on edges of cost} \leq x\},$$

261 or in the terminology of Lemma 3.3,  $\theta_0 = 2n^q\alpha$ . We show that this is sufficient for correctness  
 262 in Lemma 3.4. Choosing this type of  $\theta_0$  is important for turning the cost-scaling algorithm  
 263 into a multiplicative approximation — explained in Section B.2 — where we also give the  
 264 precise method for choosing  $\alpha$ . Note that the minimum possible  $x$  is the *minimum bottleneck*  
 265 *value* of the network.

266 Within each scale, the algorithm performs the following:

- 267 ■ SCALE-INIT takes the previous circulation (now  $2\theta$ -optimal) and transforms it into an  
 268  $\theta$ -optimal pseudoflow with  $O(k)$  excess. (For the first scale when  $\theta = \theta_0$ , SCALE-INIT  
 269 does nothing.)
- 270 ■ REFINE then reduces the excess to zero while maintaining  $\theta$ -optimality, turning  $f$  into a  
 271  $\theta$ -optimal circulation.

246 <sup>2</sup> When the costs are integers, a  $\theta$ -optimal circulation for a sufficiently small  $\theta$  (say less than  $1/n$ ) is itself  
 247 an optimal solution [13, 14]. We present this algorithm without the integral-cost assumption because in  
 248 the geometric partial matching setting (with respect to  $L_p$  norms) the costs are generally not integers.

272 If  $\theta \leq \theta^*$ , then  $f$  is a  $\theta^*$ -optimal circulation and we are done. Otherwise,  $\theta$  is halved and  
 273 the next scale begins. The algorithm produces a  $\theta^*$ -optimal circulation after  $O(\log(\theta_0/\theta^*))$   
 274 scales. Using the reduction in Lemma 3.3, we have an initial scale  $\theta_0 = 2n^q\alpha$  and target  
 275 scale  $\theta^* = \varepsilon\alpha/6k$ . Thus, the number of cost scales after the reduction is  $O(\log(n/\varepsilon))$ .  
 276  $\llbracket \log(kn/\varepsilon) \rrbracket$

277 **Scale initialization.** For SCALE-INIT, we use a simple extension to a scale initialization  
 278 procedure proposed for partial matching [13, Section 6.1]. Namely, the potential changes are  
 279 the same, but we take the additional step to remove arcs violating  $\theta$ -optimality. The vertex  
 280 set of  $H$  consists of two point sets  $A$  and  $B$ , as well as two dummy vertices  $s$  and  $t$ . The  
 281 directed edges in  $H$  are pointed from  $s$  to  $A$ , from  $A$  to  $B$ , and from  $B$  to  $t$ . We call those  
 282 arcs in  $N_H$  whose direction is consistent with their corresponding directed edges as *forward*  
 283 *arcs*, and those arcs that points in the opposite direction as *backward arcs*.

284 The procedure SCALE-INIT transforms a  $2\theta$ -optimal circulation from the previous cost  
 285 scale into a  $\theta$ -optimal flow with  $O(k)$  excess, by raising the potentials  $\pi$  of all vertices in  $A$   
 286 by  $\theta$ , those in  $B$  by  $2\theta$ , and the potential of  $t$  by  $3\theta$ . The potential of  $s$  remains unchanged.  
 287 Now the reduced cost of every forward arc is increased by  $\theta$ , and thus all the forward arcs  
 288 have reduced cost at least  $-\theta$ .

289 As for backward arcs, the procedure SCALE-INIT continues by setting the flow on  $v \rightarrow w$  to  
 290 zero for each backward arc  $w \rightarrow v$  violating the  $\theta$ -optimality constraint. In other words, we  
 291 set  $f(v \rightarrow w) = 0$  whenever  $c_\pi(w \rightarrow v) < -\theta$ . This ensures that all such backward arcs are no  
 292 longer residual, and therefore the flow (now with excess) is  $\theta$ -optimal.

293 Because the arcs are of unit-capacity in  $N_H$ , each arc desaturation creates one unit of  
 294 excess. By Lemma 3.1 the number of backward arcs is at most  $3k$ . Thus the total amount of  
 295 excess created is also  $O(k)$ . In total, the whole procedure SCALE-INIT takes  $O(n)$  time.

296 **Refinement.** The procedure REFINE is implemented using a primal-dual augmentation  
 297 algorithm, which sends flows on admissible arcs to reduce the total excess. Unlike the  
 298 Hungarian algorithm, it uses *blocking flows* instead of augmenting paths. We call a pseudoflow  
 299  $f$  on residual network  $N_g$  a *blocking flow* if  $f$  saturates at least one residual arc in every  
 300 augmenting path in  $N_g$ . In other words, there is no admissible augmenting path in  $N_{f+g}$   
 301 from an excess vertex to a deficit vertex.

302 Each iteration of REFINE finds an admissible blocking flow that is then added to the  
 303 current pseudoflow in two stages:

- 304 1. A *Hungarian search*, which increases the dual variables  $\pi$  of vertices that are reachable  
 305 from an excess vertex by at least  $\theta$ , in a Dijkstra-like manner, until there is an excess-deficit  
 306 path of admissible edges.
- 307 2. A *depth-first search* through the set of admissible arcs to construct a blocking flow. It  
 308 suffices to repeatedly extract admissible augmenting paths until no more admissible  
 309 excess-deficit paths remain.

310 The algorithm continues until the total excess becomes zero.

311 First we analyze the number of iterations executed by REFINE. The proof follows the  
 312 strategy in Goldberg *et al.* [13, Section 3.2], which proves a  $O(\sqrt{n})$  bound on the number of  
 313 iterations for general unit-capacity networks. They claimed an analogous  $O(\sqrt{k})$  bound for  
 314 partial matchings in [13, Section 6.1], which we reproduce for completeness. Due to space  
 315 constraint we have moved the proofs to Appendix B.3.

316 In the following lemma, we prove the number of iterations of REFINE required at each  
 317 cost scale. This is also the (only) place where our assumption on  $\theta_0$  become important.  $\llbracket \text{the}$



318 **following?**» We require that  $\theta_0$  is sufficiently high such that a circulation with bottleneck  
319 value at most  $\theta_0/2$  exists in  $N_H$ .

320 ► **Lemma 3.4.** *Suppose that  $\theta_0$  is chosen so that a circulation in  $N_H$  with bottleneck value at*  
321 *most  $\theta_0/2$  exists. At any scale, let  $f$  be a pseudoflow in  $N_H$  with  $O(k)$  excess. The procedure*  
322 *REFINE runs for  $O(\sqrt{k})$  iterations before the excess of  $f$  becomes zero.*

## 323 4 Fast Implementation of Refinement

324 The goal of the section is to show that after  $O(n \text{ polylog } n)$  time preprocessing, each Hungarian  
325 search and depth-first search can be implemented in  $O(k \text{ polylog } n)$  time. Combined with  
326 Lemma 3.4 the procedure REFINE can be implemented in  $O((n + k\sqrt{k}) \text{ polylog } n)$  time.  
327 Together with our analysis on scale initialization and the bound on number of cost scales,  
328 this concludes the proof to Theorem 1.2.

329 The Hungarian search and depth-first search are similar, traversing through the residual  
330 graph using admissible arcs starting from the excess vertices. Each step of the search  
331 procedures *relaxes* a minimum-reduced-cost arc from the set of visited vertices to an unvisited  
332 vertex, until a deficit vertex is reached. At a high level, our analysis strategy is to charge the  
333 relaxation events to the support arcs of  $f$ .

### 334 4.1 Null vertices and shortcut graph

335 As it turns out, there are some vertices visited by a relaxation event which we cannot charge  
336 to  $\text{supp}(f)$ . Unfortunately the number of such vertices can be as large as  $\Omega(n)$ . To overcome  
337 this issue, we replace the residual graph with an equivalent graph that excludes these *null*  
338 *vertices*, and run the Hungarian search and depth-first search on the resulting graph instead.

339 **Null vertices.** We say a vertex  $v$  in the residual graph  $N_f$  is a *null vertex* if  $\phi_f(v) = 0$  and  
340 no arc of  $\text{supp}(f)$  is incident to  $v$ . We use  $A_\emptyset$  and  $B_\emptyset$  to denote the null vertices  $A$  and  $B$   
341 respectively. Vertices that are not null are called *normal vertices*. A *null 2-path* is a length-2  
342 path in  $N_f$  from a normal vertex to another normal vertex, passing through a null vertex.  
343 As every vertex in  $A$  has in-degree 1 and every vertex in  $B$  has out-degree 1 in the residual  
344 graph, the null 2-paths must be of the form either  $(s, v, b)$  for some vertex  $b$  in  $B \setminus B_\emptyset$  or  
345  $(a, v, t)$  for some vertex  $a$  in  $A \setminus A_\emptyset$ . In either case, we say that the null 2-path *passes through*  
346 null vertex  $v$ . Similarly, we define the length-3 path from  $s$  to  $t$  that passes through two null  
347 vertices to be a *null 3-path*. Because reduced costs telescope for residual paths, the reduced  
348 cost of any null path does not depend on the potentials of the null vertices it passes through.

349 **Shortcut graph.** We construct the *shortcut graph*  $\tilde{H}_f$  from the reduction graph  $H$  by  
350 removing all null vertices and their incident edges, followed by inserting an arc from the  
351 head of each each null path  $\Pi$  to its tail, with cost equals to the sum of costs on the arcs.  
352 We call this arc the *shortcut* of null path  $\Pi$ , denoted as  $\text{short}(\Pi)$ . The resulting multigraph  
353  $\tilde{H}_f$  contains only normal vertices of  $H_f$ , and the reduced cost of any path between normal  
354 vertices are preserved. We argue now that  $\tilde{H}_f$  is fine as a surrogate for  $H_f$ . Let  $\tilde{\pi}$  be a  
355  $\theta$ -optimal potential on  $\tilde{H}_f$ . Construct potentials  $\pi$  on  $H_f$  which extends  $\tilde{\pi}$  to null vertices,  
356 by setting  $\pi(a) := \tilde{\pi}(s)$  for  $a \in A_\emptyset$  and  $\pi(b) := \tilde{\pi}(t)$  for  $b \in B_\emptyset$ .

357 ► **Lemma 4.1.** *Consider  $\tilde{\pi}$  a  $\theta$ -optimal potential on  $\tilde{H}_f$  and  $\pi$  the corresponding potential*  
358 *constructed on  $H_f$ . Then, (1) potential  $\pi$  is  $\theta$ -optimal on  $H_f$ , and (2) if arc  $\text{short}(\Pi)$  is*  
359 *admissible under  $\tilde{\pi}$ , then every arc in  $\Pi$  is admissible under  $\pi$ .*

## 4.2 Dynamic data structures for search procedures

**Hungarian search.** Conceptually, we execute the Hungarian search on the shortcut graph  $\tilde{H}_f$ . We describe how we can query the minimum-reduced-cost arc leaving  $\tilde{S}$  in  $O(\text{polylog } n)$  time for  $\tilde{H}_f$ , without constructing it explicitly. For this purpose, let  $S$  be a set of “reached” vertices maintained, identical to  $\tilde{S}$  except whenever a shortcut is relaxed, we add the null vertices passed by the corresponding null path to  $S$  in addition to its (normal) endpoints. Observe that the arcs of  $\tilde{H}_f$  leaving  $\tilde{S}$  fall into  $O(1)$  categories:

- non-shortcut backward arcs  $(v, w)$  with  $(w, v) \in \text{supp}(f)$ ;
- non-shortcut  $A$ -to- $B$  forward arcs;
- non-shortcut forward arcs from  $s$ -to- $A$  and from  $B$ -to- $t$ ;
- shortcut arcs  $(s, b)$  corresponding to null 2-paths from  $s$  to  $b \in (B \setminus B_\emptyset) \setminus S$ ;
- shortcut arcs  $(a, t)$  corresponding to null 2-paths from  $a \in (A \setminus A_\emptyset) \cap S$  to  $t$ ; and
- shortcut arcs  $(s, t)$  corresponding to null 3-paths.

For each category of arcs we maintain a proper data structure (either heap or BCP) to answer to the min-cost arc query.

**Depth-first search.** Depth-first search is similar to Hungarian search in that it uses the relaxation of minimum-reduced-cost arcs/null paths, this time to identify admissible arcs/null paths in a depth-first manner. Similar to the Hungarian search, for each category of arcs in  $\tilde{H}_f$  leaving  $\tilde{S}$ , we maintain a proper data structure to answer the minimum-reduced cost arc leaving a *fixed* vertex in  $\tilde{S}$  given by the query. Thus unlike Hungarian search which uses BCP data structures, we use dynamic nearest-neighbor data structures instead [16]. Each data structure requires  $O(1)$  queries and updates per relaxation. So in collaboration each relaxation can be implemented in  $O(\text{polylog } n)$  time [16].

**Time analysis.** The complete time analysis can be found in Appendix C.2, C.3, and C.4; here we sketch the ideas. First we show (in Appendix C.2) that both Hungarian search and depth-first search performs  $O(k)$  relaxations before a deficit vertex is reached, by looking at shortcut and non-shortcut relaxations separately. Both types of relaxations are eventually charged to the support size of  $f$ . As for the time analysis (see Appendix C.3), using the same rewinding mechanism as in Section 2.1, the running time of the Hungarian search and depth-first search, other than the potential updates, can be charged to the number of relaxations. Again using the trick by Vaidya [26] we can charge the potential updates of normal vertices to the number of relaxations in the Hungarian search. We never explicitly maintain the potentials on the null vertices; instead they are reconstructed whenever needed, either at the end of each iteration of refinement or when an augmentation sends flow through a null vertex. We show that such updates does not happen often in Appendix C.4. This completes the time analysis, which we summarize as follows:

► **Lemma 4.2.** *After  $O(n \text{ polylog } n)$ -time preprocessing, each Hungarian search and depth-first search can be implemented in  $O(k \text{ polylog } n)$  time.*

## 5 Unbalanced Transportation

In this section, we give an exact algorithm which solves the planar transportation problem in  $O((r^2\sqrt{n} + rn^{3/2}) \text{ polylog } n)$  time, proving Theorem 1.3. Our strategy is to use the standard reduction to the uncapacitated min-cost flow problem, and provide a fast implementation

under the geometric setting for the uncapacitated min-cost flow algorithm by Orlin [20], combined with some of the tools developed in Sections 2 and 3.

For lack of space, we only sketch Orlin's strongly polynomial-time algorithm for uncapacitated min-cost flow problem [20]. See Appendix D.1 for a brief introduction, as well as the original paper. In short, Orlin's algorithm follows the *excess-scaling* paradigm under the primal-dual framework: Maintain a *scale parameter*  $\Delta$ , initially set to  $U$ . A vertex  $v$  is *active* if  $|\phi_f(v)| \geq \alpha\Delta$  for a fixed parameter  $\alpha \in (0.5, 1)$ . Repeatedly run a *Hungarian search* that raises potentials (while maintaining dual feasibility) to create an admissible augmenting excess-deficit path between active vertices, on which we perform flow augmentations. Once there are no more active excess or deficit vertices,  $\Delta$  is halved. Each sequence of augmentations where  $\Delta$  holds a constant value is called an *excess scale*. On top of that, the algorithm performs contraction on arcs with flow value at least  $3n\Delta$  at the beginning of a scale, in which case the flow and potentials are no longer tracked, as well as aggressive  $\Delta$ -lowering under circumstances.

► **Theorem 5.1 (Orlin [20, Theorems 2 and 3]).** *Orlin's algorithm finds a set of optimal potentials after  $O(n \log n)$  scaling phases and  $O(n \log n)$  total augmentations.*

The remainder of the section focuses on showing that each augmentation can be implemented in  $O((r^2/\sqrt{n} + r\sqrt{n}) \text{polylog } n)$  time (after preprocessing). A subtle issue is that our geometric data structures must deal with real points in the plane instead of the contracted components. A solution is provided by Agarwal *et al.* [2]; for the sake of completeness, we describe the method for maintaining contractions under dynamic data structures in Appendix D.2.

**Recovering optimal flow.** Using a strategy from Agarwal *et al.* [2], we can recover the optimal flow in time  $O(n \text{polylog } n)$ . If furthermore the cost function is just the  $p$ -norm (without the  $q$ th-power), we show a structural result which is interesting on its own right: the set of admissible arcs under an optimal potential forms a planar graph. Thus, we can extract the admissible network in  $O(n \text{polylog } n)$  time thanks to the sparsity of planar graphs, and compute the flow using planar multiple-source multiple-sink maximum-flow algorithm by Borradaile *et al.* [7] which runs in  $O(n \log^3 n)$  time. For details see Appendix D.3 and D.4.

## 5.1 Support stars

To find an augmenting path, we again use a Hungarian search with geometric data structures to perform relaxations quickly. Our strategy is summarized as follows:

- Discard vertices that lead to dead ends in the search (not on a path to a deficit vertex).
- Cluster parts of the flow support, such that the number of support arcs outside clusters is  $O(r)$ . The number of relaxations we perform is proportional to the number of support arcs outside of clusters.

Querying/updating clusters degrades our amortized time per relaxation from  $O(\text{polylog } n)$  to  $O(\sqrt{n} \text{polylog } n)$ . Thus overall each augmentation takes  $O(r\sqrt{n} \text{polylog } n)$  time.

**Support stars.** The vertices of  $B$  with support degree 1 are partitioned into subsets  $\Sigma_a \subset B$  by the  $a \in A$  lying on the other end of their single support arc. We call  $\Sigma_a$  the *support star* centered at  $a \in A$ .

Roughly speaking, we would like to handle each support star as a single unit. When the Hungarian search reaches  $a$  or any  $b \in \Sigma_a$ , the entirety of  $\Sigma_a$  (as well as  $a$ ) is also admissibly-reachable and can be included into  $S$  without further potential updates. Additionally, the

only outgoing residual arcs of every  $b \in \Sigma_a$  lead to  $a$ , thus the only way to leave  $\Sigma_a \cup \{a\}$  is through an arc leaving  $a$ . Once a relaxation step reaches some  $b \in \Sigma_a$  or  $a$  itself, we would like to quickly update the state such that the rest of  $b \in \Sigma_a$  is also reached without performing relaxation steps to each individual  $b \in \Sigma_a$ .

## 5.2 Implementation details

Before describing our workaround for support stars, we analyze the number of relaxation steps for arcs outside of support stars. To this end we need to strip of some *dead* vertices—having no incident flow support edges and not an active excess or deficit vertex—that does not affect the search. We use  $A_\ell$  and  $B_\ell$  to denote vertices of points in  $A$  and  $B$  that are not dead. The details for handling such vertices can be found in Appendix D.5. For a proof of the following lemma, see Appendix D.6.

► **Lemma 5.2.** *Suppose we have stripped the graph of dead vertices. The number of relaxation steps in a Hungarian search outside of support stars is  $O(r)$ .*

**Relaxations outside support stars.** For relaxations that don't involve support star vertices, we can once again maintain a BCP data structure to query the minimum  $A_\ell$ -to- $B_\ell$  arc. To elaborate, this is the BCP between  $P = A_\ell \cap S$  and  $Q = (B_\ell \setminus (\bigcup_{a \in A_\ell} \Sigma_a)) \setminus S$ , weighted by potentials. Since the query is outside the support stars, there is at most one update per relaxation. Backward (support) arcs are kept admissible by the invariant, so we relax them immediately when they arrive at the frontier.

**Relaxing support stars.** We classify support stars into two categories: *big stars* with  $|\Sigma_a| > \sqrt{n}$ , and *small stars* with  $|\Sigma_a| \leq \sqrt{n}$ . Let  $A_{big} \subseteq A$  denote the centers of big stars and  $A_{small} \subseteq A$  denote the centers of small stars.

- For each big star  $\Sigma_a$ , we use a data structure  $\mathcal{D}_{big}(a)$  to maintain BCP between  $P = A_\ell \cap S$  and  $Q = \Sigma_a$ . We query this until  $a \in S$  or any vertex of  $\Sigma_a$  is added to  $S$ .
- All small stars are added to a single BCP data structure  $\mathcal{D}_{small}$  between  $P = A_\ell \cap S$  and  $Q = (\bigcup_{a \in A_{small}} \Sigma_a) \setminus S$ . When an  $a \in A_{small}$  or any vertex of its support star is added to  $S$ , we remove the points of  $\Sigma_a$  from  $\mathcal{D}_{small}$  using  $|\Sigma_a|$  deletion operations.

We will update these data structures as each support star center is added into  $S$ . If a relaxation step adds some  $b \in B_\ell$  and  $b$  is in a support star  $\Sigma_a$ , then we immediately relax  $b \rightarrow a$ , as all support arcs are admissible.

Suppose a relaxation step adds  $a \in A_\ell$  to  $S$ . We must (i) remove  $a$  from every  $\mathcal{D}_{big}$ , (ii) remove  $a$  from  $\mathcal{D}_{small}$ . If  $a \in A_{big}$ , we also (iii) deactivate  $\mathcal{D}_{big}(a)$ . If  $a \in A_{small}$ , we also (iv) remove the points of  $\Sigma_a$  from  $\mathcal{D}_{small}$ . The operations (i–iii) can be performed in  $O(\text{polylog } n)$  time each, but (iv) may take up to  $O(\sqrt{n} \text{polylog } n)$  time. On the other hand, there are now  $O(\sqrt{n})$  data structures to query during each relaxation step, which takes  $O(\sqrt{n} \log^2 n)$  time in total. Together with Lemma 5.2 we bound the time for each Hungarian search.

► **Lemma 5.3.** *Hungarian search takes  $O(r\sqrt{n} \text{polylog } n)$  time.*

**Updating support stars.** As the flow support changes, the membership of support stars may shift causing a big star to become small or vice versa. To efficiently support this, we introduce a soft boundary in determining whether a support star is big or small. Standard charging argument shows that the amortized update time per membership change is  $O(\text{polylog } n)$  for big-to-small updates, and  $O((r/\sqrt{n}) \text{polylog } n)$  for small-to-big ones; see Appendix D.7. The

number of star membership changes by adding a single augmenting path is bounded above by twice of its length, so  $O(r)$ . By Theorem 5.1, the total number of membership changes is  $O(rn \log n)$ . The total time spent on big-to-small updates is  $O(rn \text{polylog } n)$ , and the total time spent on small-to-big updates is  $O(r^2 \sqrt{n} \text{polylog } n)$ . Membership changes themselves can be performed in  $O(\text{polylog } n)$  time each.

**Preprocessing time.** It takes  $O(rn \text{polylog } n)$  time to build the very first set of data structures. There are  $r \cdot |\Sigma_a|$  points to insert for each  $\mathcal{D}_{\text{big}}(a)$ , so the number of points to insert is  $O(rn)$ . At most  $O(rn)$  points have to be inserted for  $\mathcal{D}_{\text{small}}$ . So the total preprocessing time is  $O(rn \text{polylog } n)$ .

**Between searches.** After each augmentation, we reset the data structures to their initial state plus the change from augmentation using rewinding mechanism (see Section 2.1). This takes time proportional to the time for Hungarian search, which is  $O(r \sqrt{n} \text{polylog } n)$  by Lemma 5.3. The most recent augmentation may have deactivated  $O(1)$  active excess and deficit vertices, which takes  $O(\sqrt{n} \text{polylog } n)$  time to update. Finally, we note that an augmenting path cannot reduce the support degree of a vertex to zero, and therefore no new dead vertices are created by augmentation.

**Between excess scales.** When the excess scale changes, vertices that were previously inactive may become active, and vertices that were dead may be revived. If we have the data structures built at the end of the previous scale, then we can add in each new active vertex  $a \in A$  and charge the insertion to the (future) augmenting path or contraction which eventually causes the vertex being inactive or absorbed. By Theorem 5.1, there are  $O(n \log n)$  such newly active vertices; each of them takes  $O(\sqrt{n} \text{polylog } n)$  to update. So the total time spent is  $O(n^{3/2} \text{polylog } n)$ .

**Putting it together.** After  $O(rn \text{polylog } n)$  preprocessing, we spend  $O(r \sqrt{n} \text{polylog } n)$  time on relaxations each Hungarian search by Lemma 5.3, for a total of  $O(rn^{3/2} \text{polylog } n)$  time over the course of the algorithm. Rewinding takes the same amount of time. We spend up to  $O((r^2 \sqrt{n} + rn) \text{polylog } n)$  time switching stars between big/small. We spend  $O(n^{3/2} \text{polylog } n)$  time activating and reviving vertices. Adding, the algorithm takes  $O((r^2 \sqrt{n} + rn^{3/2}) \text{polylog } n)$  time to produce optimal potentials  $\pi^*$ , from which we can recover  $f^*$  in  $O(n \text{polylog } n)$  additional time. This completes the proof of Theorem 1.3.

## References

- 1 Pankaj K. Agarwal, Alon Efrat, and Micha Sharir. Vertical decomposition of shallow levels in 3-dimensional arrangements and its applications. *SIAM J. Comput.* 29(3):912–953, 1999. <https://doi.org/10.1137/S0097539795295936>.
- 2 Pankaj K. Agarwal, Kyle Fox, Debmalya Panigrahi, Kasturi R. Varadarajan, and Allen Xiao. Faster algorithms for the geometric transportation problem. *33rd International Symposium on Computational Geometry, SoCG 2017, July 4-7, 2017, Brisbane, Australia*, 7:1–7:16, 2017. <https://doi.org/10.4230/LIPIcs.SoCG.2017.7>.
- 3 Pankaj K. Agarwal, Haim Kaplan, Geva Kipper, Wolfgang Mulzer, Günter Rote, Micha Sharir, and Allen Xiao. Approximate minimum-weight matching with outliers under translation. *CoRR* abs/1810.10466, 2018. <http://arxiv.org/abs/1810.10466>.
- 4 Pankaj K. Agarwal and Kasturi R. Varadarajan. A near-linear constant-factor approximation for euclidean bipartite matching? *Proceedings of the 20th ACM Symposium on*

- 531      *Computational Geometry, Brooklyn, New York, USA, June 8-11, 2004*, 247–252, 2004.  
 532      <https://doi.org/10.1145/997817.997856>).
- 533      **5** Rinat Ben Avraham, Matthias Henze, Rafel Jaume, Balázs Keszegh, Orit E. Raz, Micha  
 534      Sharir, and Igor Tubis. Partial-matching RMS distance under translation: Combinator-  
 535      ics and algorithms. *Algorithmica* 80(8):2400–2421, 2018. [https://doi.org/10.1007/](https://doi.org/10.1007/s00453-017-0326-0)  
 536      [s00453-017-0326-0](https://doi.org/10.1007/s00453-017-0326-0)).
- 537      **6** D. Bertsekas and D. El Baz. Distributed asynchronous relaxation methods for convex  
 538      network flow problems. *SIAM Journal on Control and Optimization* 25(1):74–85, 1987.  
 539      <https://doi.org/10.1137/0325006>).
- 540      **7** Glencora Borradaile, Philip N. Klein, Shay Mozes, Yahav Nussbaum, and Christian Wulff-  
 541      Nilsen. Multiple-source multiple-sink maximum flow in directed planar graphs in near-  
 542      linear time. *SIAM J. Comput.* 46(4):1280–1303, 2017. [https://doi.org/10.1137/](https://doi.org/10.1137/15M1042929)  
 543      [15M1042929](https://doi.org/10.1137/15M1042929)).
- 544      **8** Sergio Cabello, Panos Giannopoulos, Christian Knauer, and Günter Rote. Matching point  
 545      sets with respect to the earth mover’s distance. *Comput. Geom.* 39(2):118–133, 2008.  
 546      <https://doi.org/10.1016/j.comgeo.2006.10.001>).
- 547      **9** Ran Duan, Seth Pettie, and Hsin-Hao Su. Scaling algorithms for approximate and exact  
 548      maximum weight matching. *CoRR* abs/1112.0790, 2011. [http://arxiv.org/abs/1112.](http://arxiv.org/abs/1112.0790)  
 549      [0790](http://arxiv.org/abs/1112.0790)).
- 550      **10** Jack Edmonds and Richard M. Karp. Theoretical improvements in algorithmic efficiency  
 551      for network flow problems. *J. ACM* 19(2):248–264, 1972. [https://doi.org/10.1145/](https://doi.org/10.1145/321694.321699)  
 552      [321694.321699](https://doi.org/10.1145/321694.321699)).
- 553      **11** Shimon Even and Robert E. Tarjan. Network flow and testing graph connectivity. *SIAM*  
 554      *J. Comput.* 4(4):507–518, 1975. <https://doi.org/10.1137/0204043>).
- 555      **12** Harold N. Gabow and Robert E. Tarjan. Faster scaling algorithms for network problems.  
 556      *SIAM J. Comput.* 18(5):1013–1036, 1989. <https://doi.org/10.1137/0218069>).
- 557      **13** Andrew V. Goldberg, Sagi Hed, Haim Kaplan, and Robert E. Tarjan. Minimum-cost  
 558      flows in unit-capacity networks. *Theory Comput. Syst.* 61(4):987–1010, 2017. <https://doi.org/10.1007/s00224-017-9776-7>).
- 559      **14** Andrew V. Goldberg and Robert E. Tarjan. Finding minimum-cost circulations by success-  
 560      ive approximation. *Math. Oper. Res.* 15(3):430–466, 1990. [https://doi.org/10.1287/](https://doi.org/10.1287/moor.15.3.430)  
 561      [moor.15.3.430](https://doi.org/10.1287/moor.15.3.430)).
- 562      **15** John E. Hopcroft and Richard M. Karp. An  $n^{5/2}$  algorithm for maximum matchings in  
 563      bipartite graphs. *SIAM J. Comput.* 2(4):225–231, 1973. [https://doi.org/10.1137/](https://doi.org/10.1137/0202019)  
 564      [0202019](https://doi.org/10.1137/0202019)).
- 565      **16** Haim Kaplan, Wolfgang Mulzer, Liam Roditty, Paul Seiferth, and Micha Sharir. Dynamic  
 566      planar Voronoi diagrams for general distance functions and their algorithmic applications.  
 567      *Proceedings of the Twenty-Eighth Annual ACM-SIAM Symposium on Discrete Algorithms,*  
 568      *SODA 2017, Barcelona, Spain, Hotel Porta Fira, January 16-19*, 2495–2504, 2017. <https://doi.org/10.1137/1.9781611974782.165>).
- 569      **17** Harold W Kuhn. The Hungarian method for the assignment problem. *Naval Research*  
 570      *Logistics (NRL)* 2(1-2):83–97. Wiley Online Library, 1955.
- 571      **18** Yin Tat Lee and Aaron Sidford. Following the path of least resistance : An  $\tilde{O}(m \sqrt{n})$   
 572      algorithm for the minimum cost flow problem. *CoRR* abs/1312.6713, 2013. [http://arxiv.](http://arxiv.org/abs/1312.6713)  
 573      [org/abs/1312.6713](http://arxiv.org/abs/1312.6713)).
- 574      **19** Aleksander Mądry. Navigating central path with electrical flows: From flows to matchings,  
 575      and back. *54th Annual IEEE Symposium on Foundations of Computer Science, FOCS*  
 576      *2013, 26-29 October, 2013, Berkeley, CA, USA*, 253–262, 2013. [https://doi.org/10.](https://doi.org/10.1109/FOCS.2013.35)  
 577      [1109/FOCS.2013.35](https://doi.org/10.1109/FOCS.2013.35)).



- 580 20 James B. Orlin. A faster strongly polynomial minimum cost flow algorithm. *Operations Research* 41(2):338–350, 1993. (<https://doi.org/10.1287/opre.41.2.338>).
- 581 21 Lyle Ramshaw and Robert E. Tarjan. A weight-scaling algorithm for min-cost imperfect matchings in bipartite graphs. *53rd Annual IEEE Symposium on Foundations of Computer Science, FOCS 2012, New Brunswick, NJ, USA, October 20-23, 2012*, 581–590, 2012. (<https://doi.org/10.1109/FOCS.2012.9>).
- 582 22 Günter Rote. Partial least-squares point matching under translations. *Proc. 26th European Workshop Comput. Geom. (EuroCG 2010)* 249–251, 2010.
- 583 23 R. Sharathkumar and Pankaj K. Agarwal. Algorithms for the transportation problem in geometric settings. *Proceedings of the Twenty-Third Annual ACM-SIAM Symposium on Discrete Algorithms, SODA 2012, Kyoto, Japan, January 17-19, 2012*, 306–317, 2012. (<http://portal.acm.org/citation.cfm?id=2095145&CFID=63838676&CFTOKEN=79617016>).
- 584 24 R. Sharathkumar and Pankaj K. Agarwal. A near-linear time  $\epsilon$ -approximation algorithm for geometric bipartite matching. *Proceedings of the 44th Symposium on Theory of Computing Conference, STOC 2012, New York, NY, USA, May 19 - 22, 2012*, 385–394, 2012. (<https://doi.org/10.1145/2213977.2214014>).
- 585 25 Éva Tardos. A strongly polynomial minimum cost circulation algorithm. *Combinatorica* 5(3):247–256, 1985. (<https://doi.org/10.1007/BF02579369>).
- 586 26 Pravin M. Vaidya. Geometry helps in matching. *SIAM J. Comput.* 18(6):1201–1225, 1989. (<https://doi.org/10.1137/0218080>).
- 587 27 Kasturi R. Varadarajan. A divide-and-conquer algorithm for min-cost perfect matching in the plane. *39th Annual Symposium on Foundations of Computer Science, FOCS '98, November 8-11, 1998, Palo Alto, California, USA*, 320–331, 1998. (<https://doi.org/10.1109/SFCS.1998.743466>).

## 604 **A** Missing Details and Proofs from Section 2

### 605 **A.1** Potential updates

606 We modify a trick from Vaidya [26] to batch potential updates. Potentials have a *stored*  
 607 *value*, i.e. the currently recorded value of  $\pi(v)$ , and a *true value*, which may have changed  
 608 from  $\pi(v)$ . The resulting algorithm queries the minimum-reduced-cost under the true values  
 609 of  $\pi$  and updates the stored value occasionally.

610 Throughout the entire Hungarian algorithm, we maintain a nonnegative scalar  $\delta$  (initially  
 611 set to 0) which aggregates potential changes. Vertices  $a \in A$  that are added to  $S$  are inserted  
 612 into BCP with weight  $\omega(a) \leftarrow \pi(a) - \delta$ , for whatever value  $\delta$  is at the time of insertion.  
 613 Similarly, vertices  $b \in B$  that are added to  $S$  have  $\omega(b) \leftarrow \pi(b) - \delta$  recorded ( $B \cap S$  points  
 614 aren't added into a BCP set). When the Hungarian search wants to raise the potentials of  
 615 points in  $S$ ,  $\delta$  is increased by that amount instead. Thus, true value for any potential of a  
 616 point in  $S$  is always  $\omega(p) + \delta$ . For points of  $(A \cup B) \setminus S$ , the true potential is equal to the  
 617 stored potential. Since all the points of  $A \cap S$  have weights uniformly offset from their true  
 618 potentials, the minimum edge returned by the BCP does not change. *«why?»*

619 Once a point is removed from  $S$  (i.e. by an augmentation or the rewinding mechanism),  
 620 we update its stored potential  $\pi(p) \leftarrow \omega(p) + \delta$ , again for the current value of  $\delta$ . Most  
 621 importantly,  $\delta$  is not reset at the end of a Hungarian search and persists through the entire  
 622 algorithm. Thus, the initial BCP sets constructed by the rewinding mechanism have true  
 623 potentials accurately represented by  $\delta$  and  $\omega(p)$ .

624 We update  $\delta$  once per edge relaxations; thus  $O(k)$  times in total per Hungarian search.  
 625 There are  $O(k)$  stored values updated per Hungarian search during the rewinding process.  
 626 The time spent on potential updates per Hungarian search is therefore  $O(k)$ .

## 627 **B** Missing Details and Proofs from Section 3

### 628 **B.1** Preliminaries on Network Flows

629 **Network.** Let  $G = (V, E)$  be a directed graph, augmented by edge costs  $c$  and capacities  
 630  $u$ , and a supply-demand function  $\phi$  defined on the vertices. We can turn the graph  $G$  into  
 631 a *network*  $N = (V, \vec{E})$ : For each directed edge  $(v, w)$  in  $E$ , insert two *arcs*  $v \rightarrow w$  and  $w \rightarrow v$   
 632 into the arc set  $\vec{E}$ ; the *forward arc*  $v \rightarrow w$  inherits the capacity and cost from the directed  
 633 graph  $G$ , while the *backward arc*  $w \rightarrow v$  satisfies  $u(w \rightarrow v) = 0$  and  $c(w \rightarrow v) = -c(v \rightarrow w)$ . This  
 634 we ensure that the graph  $(V, \vec{E})$  is *symmetric* and the cost function  $c$  is *antisymmetric* on  
 635  $N$ . The positive values of  $\phi(v)$  are referred to as *supply*, and the negative values of  $\phi(v)$   
 636 as *demand*. We assume that all capacities are nonnegative, all supplies and demands are  
 637 integers, and the sum of supplies and demands is equal to zero. A *unit-capacity* network has  
 638 all its edge capacities equal to 1. In this section we assume all networks are of unit-capacity.

641 **Pseudoflows.** Given a network  $N := (V, \vec{E}, c, u, \phi)$ , a *pseudoflow* (or *flow* to be short)  
 642  $f: \vec{E} \rightarrow \mathbb{Z}$  on  $N$  is an antisymmetric function on the arcs of  $N$  satisfying  $f(v \rightarrow w) \leq u(v \rightarrow w)$   
 643 for every arc  $v \rightarrow w$ .<sup>3</sup> We say that  $f$  *saturates* an arc  $v \rightarrow w$  if  $f(v \rightarrow w) = u(v \rightarrow w)$ ; an arc  $v \rightarrow w$   
 644 is *residual* if  $f(v \rightarrow w) < u(v \rightarrow w)$ . The *support* of  $f$  in  $N$ , denoted as  $\text{supp}(f)$ , is the set of

---

639 <sup>3</sup> In general the pseudoflows are allowed to take real-values. Here under the unit-capacity assumption any  
 640 optimal flows are integer-valued.

arcs with positive flows:

$$\text{supp}(f) := \{v \rightarrow w \in \vec{E} \mid f(v \rightarrow w) > 0\}.$$

Given a pseudoflow  $f$ , we define the *imbalance* of a vertex (with respect to  $f$ ) to be

$$\phi_f(v) := \phi(v) + \sum_{w \rightarrow v \in \vec{E}} f(w \rightarrow v) - \sum_{v \rightarrow w \in \vec{E}} f(v \rightarrow w).$$

We call positive imbalance *excess* and negative imbalance *deficit*; and vertices with positive (resp. negative) imbalance *excess* (resp. *deficit*) *vertices*. A vertex is *balanced* if it has zero imbalance. If all vertices are balanced, the pseudoflow is a *circulation*. The *cost* of a pseudoflow is defined to be

$$\text{cost}(f) := \sum_{v \rightarrow w \in \text{supp}(f)} c(v \rightarrow w) \cdot f(v \rightarrow w).$$

The *minimum-cost flow problem (MCF)* asks to find a circulation of minimum cost inside a given network.

**Residual graph.** Given a pseudoflow  $f$ , the *residual network* is defined as follows. Recall that the set of *residual arcs*  $\vec{E}_f$  under  $f$  are those arcs  $v \rightarrow w$  satisfying  $f(v \rightarrow w) < u(v \rightarrow w)$ . In other words, an arc that is not saturated by  $f$  is a residual arc; similarly, given an arc  $v \rightarrow w$  with positive flow value, the backward arc  $w \rightarrow v$  is a residual arc.

Let  $N = (V, \vec{E}, c, u, \phi)$  be a network constructed from graph  $G$ , with a pseudoflow  $f$  on  $N$ . The *residual graph*  $G_f$  of  $f$  has  $V$  as its vertex set and  $\vec{E}_f$  as its arc set. The *residual capacity*  $u_f$  with respect to pseudoflow  $f$  is defined to be  $u_f(v \rightarrow w) := u(v \rightarrow w) - f(v \rightarrow w)$ . Observe that the residual capacity is always nonnegative. We can define residual arcs differently using residual capacities:

$$\vec{E}_f = \{v \rightarrow w \mid u_f(v \rightarrow w) > 0\}.$$

In other words, the set of residual arcs are precisely those arcs in the residual graph, each of which has nonzero residual capacity.

#### Admissible flow augmentation.

► **Lemma B.1.** *Let  $f$  be an  $\varepsilon$ -optimal pseudoflow in  $G$  and let  $f'$  be an admissible flow in  $G_f$ . Then  $f + f'$  is also  $\theta$ -optimal.*

**Proof.** Augmentation by  $f'$  will not change the potentials, so any previously  $\theta$ -optimal arcs remain  $\theta$ -optimal. However, it may introduce new arcs  $v \rightarrow w$  with  $u_{f+f'}(v \rightarrow w) > 0$ , that previously had  $u_f(v \rightarrow w) = 0$ . We will verify that these arcs satisfy the  $\theta$ -optimality condition.

If an arc  $v \rightarrow w$  is newly introduced this way, then by definition of residual capacities  $f(v \rightarrow w) = u(v \rightarrow w)$ . At the same time,  $u_{f+f'}(v \rightarrow w) > 0$  implies that  $(f + f')(v \rightarrow w) < u(v \rightarrow w)$ . This means that  $f'$  augmented flow in the reverse direction of  $v \rightarrow w$ , that is,  $f'(w \rightarrow v) > 0$ . By assumption, the arcs of  $\text{supp}(f')$  are admissible, so  $w \rightarrow v$  was an admissible arc ( $c_\pi(w \rightarrow v) \leq 0$ ). By antisymmetry of reduced costs, this implies  $c_\pi(v \rightarrow w) \geq 0 \geq -\theta$ . Therefore, all arcs with  $u_{f+f'}(v, w) > 0$  respect the  $\theta$ -optimality condition, and thus  $f + f'$  is  $\theta$ -optimal. ◀

► **Lemma 3.1.** *The size of  $\text{supp}(f)$  is at most  $3k$  for any integer circulation  $f$  in reduction network  $N_H$ . As a corollary, the number of residual backward arcs is at most  $3k$ .*

**Proof.** Because  $f$  is a circulation,  $\text{supp}(f)$  can be decomposed into  $k$  paths from  $s$  to  $t$ . Each  $s$ -to- $t$  path in  $N_H$  is of length three, so the size of  $\text{supp}(f)$  is at most  $3k$ . As every backward arc in the residual network must be induced by positive flow in the opposite direction, the total number of residual backward arcs is at most  $3k$ . ◀

► **Lemma 3.2.** *Let  $f$  be a  $\theta$ -optimal integer circulation in  $N_H$ , and  $f^*$  be an optimal integer circulation for  $N_H$ . Then,  $\text{cost}(f) \leq \text{cost}(f^*) + 6k\theta$ .*

**Proof.** By Lemma 3.1, the total number of backward arcs in the residual network  $N_f$  is at most  $3k$ . Consider the residual flow in  $N_f$  defined by the difference between  $f^*$  and  $f$ . Since both  $f$  and  $f^*$  are both circulations and  $N_H$  has unit-capacity, the flow  $f - f^*$  is comprised of unit flows on a collection of edge-disjoint residual cycles  $\Gamma_1, \dots, \Gamma_\ell$ . Observe that each residual cycle  $\Gamma_i$  must have exactly half of its arcs being backward arcs, and thus we have  $\sum_i |\Gamma_i| \leq 6k$ .

Let  $\pi$  be some potential certifying that  $f$  is  $\theta$ -optimal. Because  $\Gamma_i$  is a residual cycle, we have  $c_\pi(\Gamma_i) = c(\Gamma_i)$  since the potential terms telescope. We then see that

$$\text{cost}(f) - \text{cost}(f^*) = \sum_i c(\Gamma_i) = \sum_i c_\pi(\Gamma_i) \geq \sum_i (-\theta) \cdot |\Gamma_i| \geq -6k\theta,$$

where the second-to-last inequality follows from the  $\theta$ -optimality of  $f$  with respect to  $\pi$ . Rearranging the terms we have that  $\text{cost}(f) \leq \text{cost}(f^*) + 6k\theta$ . ◀

## B.2 Multiplicative approximation

◀ **«To be compressed.»** The reduction in Sharathkumar and Agarwal [23] uses a *single-linkage clustering* of  $G$ —the clustering induced by a low-cost subgraph of the minimum spanning tree (MST) on  $G$ —to transform the instance into a number of bounded-diameter point sets “preserving”  $M^*$ . We also perform a single-linkage clustering, but use it to find a good upper bound on the minimum bottleneck value of the reduction network instead.

Let  $T$  be the minimum spanning tree on input graph  $G$  and order its edges by increasing length as  $e_1, \dots, e_{r+n-1}$ ; both  $T$  and the sorted edges can be constructed in  $O(n \text{ polylog } n)$  time using a dynamic data structure for planar nearest neighbors. **«HC: Don’t mix construction with analysis.»** Let  $T_i$  denote the connected components of  $T$  induced by the first  $i$  edges  $(e_1, \dots, e_i)$  of  $T$ . Abusing notation, we also use  $T_i$  to denote the clustering of  $A \cup B$  induced by the components of  $T_i$ . For each cluster  $K \in T_i$ , we use  $A_K$  and  $B_K$  to denote the points of  $A \cap K$  and  $B \cap K$  respectively.

We say a partial matching  $M$  is *intra-cluster* with respect to a clustering of  $A \cup B$  if no edge of  $M$  has endpoints in two different clusters. Otherwise, the matching is *inter-cluster*. Let  $i^*$  be the smallest index such that  $T_{i^*}$  contains an intra-cluster partial matching. We can compute  $i^*$  in  $O(n)$  time by traversing the edges  $e_1, \dots, e_{r+n-1}$  in ascending order, tracking the number of  $A$  and  $B$  points in each cluster defined by  $T_i$ . Specifically, for each  $i \in 1, \dots, (r+n-1)$  let

$$\text{count}(i) := \sum_{K \in T_i} \min(|A_K|, |B_K|),$$

which is the size of the largest intra-cluster matching possible under  $T_i$ . The first index for which  $\text{count}(i) \geq k$  is  $i^*$ .

► **Lemma B.2.**

- (i)  $c(e_{i^*}) \leq \text{cost}(M^*)$
- (ii) *The minimum bottleneck value for the reduction network on  $A \cup B$  is at most  $n^q \cdot c(e_{i^*})$ .*

**Proof.** First, by definition of  $i^*$ ,  $M^*$  is inter-cluster in  $T_{i^*-1}$ . Although  $M^*$  need not be intra-cluster in  $T_{i^*}$ , this implies that at least one edge  $e \in M^*$  has  $c(e) \geq c(e_{i^*})$ . It follows that  $c(e_{i^*}) \leq \text{cost}(M^*)$ .

Second, let  $M$  be the intra-cluster  $k$ -matching purported to exist in  $T_{i^*}$ . Each component of  $T_{i^*}$  has first-power diameter at most  $n\|e_{i^*}\|$ , so the longest edge in  $M$  has cost at most  $n^q \cdot c(e_{i^*})$ . In the reduction network  $N$ ,  $M$  corresponds to a circulation where the most expensive arc costs at most  $n^q \cdot c(e_{i^*})$ ; the minimum bottleneck value of  $N$  is at most that. ◀

With Lemma B.2, we can complete the proof of Lemma 3.3.

**Proof of Lemma 3.3.** We can determine  $i^*$  and  $c(e_{i^*})$  in  $O(n \text{ polylog } n)$  time, as described above. Fix  $\alpha = c(e_{i^*})$ . By Lemma B.2.ii the minimum bottleneck value of the reduction network is at most  $n^q \alpha$ . Computing an  $(\alpha\varepsilon/6k)$ -optimal circulation produces a partial matching of cost at most

$$\text{cost}(M^*) + 6k(\alpha\varepsilon/6k) = \text{cost}(M^*) + \varepsilon \cdot c(e_{i^*}) \leq (1 + \varepsilon) \text{cost}(M^*)$$

where the inequality follows from Lemma B.2.i. ◀

### B.3 Number of iterations during refinement.

To this end we need a bound on the size of the support of  $f$  right before and throughout the execution of REFINE. This bound will also be useful later for bounding the overall running time.

► **Lemma B.3.** *Let  $f$  be an integer pseudoflow in  $N_H$  with  $O(k)$  excess. Then, the size of the support of  $f$  is at most  $O(k)$ .*

**Proof.** Observe that the reduction graph  $H$  is a directed acyclic graph, and thus the support of  $f$  does not contain a cycle. Now  $\text{supp}(f)$  can be decomposed into a set of inclusion-maximal paths, each of which contributes a single unit of excess to the flow if the path does not terminate at  $t$  or if more than  $k$  paths terminate at  $t$ . By assumption, there are  $O(k)$  units of excess to which we can associate to the paths, and at most  $k$  paths (those that terminate at  $t$ ) that we cannot associate with a unit of excess. The length of any such paths is at most three by construction of the reduction graph  $H$ . Therefore we can conclude that the number of arcs in the support of  $f$  is  $O(k)$ . ◀

► **Corollary B.4.** *The size of  $\text{supp}(f)$  is at most  $O(k)$  for pseudoflow  $f$  right before or during the execution of REFINE.*

► **Lemma 3.4.** *Suppose that  $\theta_0$  is chosen so that a circulation in  $N_H$  with bottleneck value at most  $\theta_0/2$  exists. At any scale, let  $f$  be a pseudoflow in  $N_H$  with  $O(k)$  excess. The procedure REFINE runs for  $O(\sqrt{k})$  iterations before the excess of  $f$  becomes zero.*

**Proof.** Let  $f_0$  and  $\pi_0$  be the circulation and potential at the end of the previous cost scale. For the first cost scale, let  $\pi_0 := 0$  and  $f_0$  be any circulation with bottleneck value  $\leq \theta_0/2$ . In both cases,  $f_0$  is a  $2\theta$ -optimal circulation with respect to  $\pi_0$ . Let  $f$  and  $\pi$  be the current flow and the potential. Let  $d(v)$  defined to be the amount of potential increase at  $v$  since  $\pi_0$ , measured in units of  $\theta$ ; in other words,  $d(v) := (\pi(v) - \pi_0(v))/\theta$ .

Now divide the iterations executed by the procedure REFINE into two phases: The transition from the first phase to the second happens when every excess vertex  $v$  has  $d(v) \geq \sqrt{k}$ . At most  $\sqrt{k}$  iterations belong to the first phase as each Hungarian search

increases the potential  $\pi$  by at least  $\theta$  for each excess vertex (and thus increases  $d(v)$  by at least one).

The number of iterations belonging to the second phase is upper bounded by the amount of total excess at the end of the first phase, because each subsequent push of a blocking flow reduces the total excess by at least one. We now show that the amount of such excess is at most  $O(\sqrt{k})$ . Consider the set of arcs  $E^+ := \{v \rightarrow w \mid f(v \rightarrow w) < f_0(v \rightarrow w)\}$ . The total amount of excess is upper bounded by the number of arcs in  $E^+$  that crosses an arbitrarily given cut  $X$  that separates the excess vertices from the deficit vertices, when the network has unit-capacity [13, Lemma 3.6]. Consider the set of cuts  $X_i := \{v \mid d(v) > i\}$  for  $0 \leq i < \sqrt{k}$ ; every such cut separates the excess vertices from the deficit vertices at the end of first phase. Each arc in  $E^+$  crosses at most 3 cuts of type  $X_i$  [13, Lemma 3.1]. So there is one  $X_i$  crossed by at most  $3|E^+|/\sqrt{k}$  arcs in  $E^+$ . The size of  $E^+$  is bounded by the sum of support sizes of  $f$  and  $f_0$ ; by Corollary B.4 the size of  $E^+$  is  $O(k)$ . This implies an  $O(\sqrt{k})$  bound on the total excess after the first phase, which in turn bounds the number of iterations in the second phase. ◀

## C Missing Details and Proofs from Section 4

► **Lemma 4.1.** *Consider  $\tilde{\pi}$  a  $\theta$ -optimal potential on  $\tilde{H}_f$  and  $\pi$  the corresponding potential constructed on  $H_f$ . Then, (1) potential  $\pi$  is  $\theta$ -optimal on  $H_f$ , and (2) if arc  $\text{short}(\Pi)$  is admissible under  $\tilde{\pi}$ , then every arc in  $\Pi$  is admissible under  $\pi$ .*

**Proof.** Reduced costs for any arc from a normal vertex another is unchanged under either  $\tilde{\pi}$  or  $\pi$ . Recall that a null path is comprised of one  $A$ -to- $B$  arc, and one or two zero-cost arcs (connecting the null vertex/vertices to  $s$  and/or  $t$ ). With our choice of null vertex potentials, we observe that the zero-cost arcs still have zero reduced cost. It remains to prove that an arbitrary **«residual?»** arc  $(a, b)$  **«arc or directed edge?»** satisfies the  $\theta$ -optimality condition and admissibility when either  $a$  or  $b$  is a null vertex.

By construction of the shortcut graph, there is always a null path  $\Pi$  that contains  $(a, b)$ . Observe that  $c_\pi(a, b) = c_\pi(\Pi)$ , independent to the type of null path. Again by construction,  $c_\pi(\Pi) = c_{\tilde{\pi}}(\text{short}(\Pi))$ , so we have  $c_\pi(a, b) = c_{\tilde{\pi}}(\text{short}(\Pi)) \geq -\theta$ . Additionally, if  $\text{short}(\Pi)$  is admissible under  $\tilde{\pi}$ , then so is  $(a, b)$  under  $\pi$ . This proves the lemma. ◀

### C.1 Dynamic data structures for search procedures

Here we formally describe in details the set of dynamic data structure we use for the Hungarian search and depth-first search procedures.

For Hungarian search, we maintain the following for each type of outgoing arcs of  $\tilde{H}_f$  leaving  $\tilde{S}$ :

1. Non-shortcut backward arcs  $(v, w)$  with  $(w, v) \in \text{supp}(f)$ . For these, we can maintain a min-heap on  $\text{supp}(f)$  arcs as each  $v$  arrives in  $\tilde{S}$ .
2. Non-shortcut  $A$ -to- $B$  forward arcs. For these, we can use a BCP data structure between  $(A \setminus A_\emptyset) \cap \tilde{S}$  and  $(B \setminus B_\emptyset) \setminus \tilde{S}$ , weighted by potential.
3. Non-shortcut forward arcs from  $s$ -to- $A$  and from  $B$ -to- $t$ . For  $s$ , we can maintain a min-heap on the potentials of  $B \setminus \tilde{S}$ , queried while  $s \in \tilde{S}$ . For  $t$ , we can maintain a max-heap on the potentials of  $A \cap \tilde{S}$ , queried while  $t \notin \tilde{S}$ .
4. Shortcut arcs  $(s, b)$  corresponding to null 2-paths from  $s$  to  $b \in (B \setminus B_\emptyset) \setminus S$ . For these, we maintain a BCP data structure with  $P = A_\emptyset$ ,  $Q = (B \setminus B_\emptyset) \setminus S$  with weights  $\omega(p) = \pi(s)$



for all  $p \in P$ , and  $\omega(q) = \pi(q)$  for all  $q \in Q$ . A response  $(a, b)$  corresponds to th null 2-path  $(s, a, b)$ . This is only queried while  $s \in S$ .

5. Shortcut arcs  $(a, t)$  corresponding to null 2-paths from  $a \in (A \setminus A_\emptyset) \cap S$  to  $t$ . For these, we maintain a BCP data structure with  $P = (A \setminus A_\emptyset) \cap S$ ,  $Q = B_\emptyset \setminus S$  with weights  $\omega(p) = \pi(p)$  for all  $p \in P$ , and  $\omega(q) = \pi(t)$  for all  $q \in Q$ . A response  $(a, b)$  corresponds to th null 2-path  $(a, b, t)$ . This is only queried while  $t \notin S$ .
6. Shortcut arcs  $(s, t)$  corresponding to null 3-paths. For these, we maintain in a BCP data structure with  $P = A_\emptyset \setminus S$ ,  $Q = B_\emptyset \setminus S$  with weights  $\omega(p) = \pi(s)$  for all  $p \in P$ , and  $\omega(q) = \pi(t)$  for all  $q \in Q$ . A response  $(a, b)$  corresponds to th null 3-path  $(s, a, b, t)$ . This is only queried while  $s \in S$  and  $t \notin S$ .

For depth-first search, we maintain the following for each type of outgoing arcs of  $\tilde{H}_f$  leaving  $\tilde{S}$ :

1. Non-shortcut backward arcs  $(v', w')$  with  $(w', v') \in \text{supp}(f)$ . For these, we can maintain a min-heap on  $(w', v') \in \text{supp}(f)$  arcs for each normal  $v' \in V$ .
2. Non-shortcut  $A$ -to- $B$  forward arcs. For these, we maintain a NN data structure over  $P = (B \setminus B_\emptyset) \setminus \tilde{S}$ , with weights  $\omega(p) = \pi(p)$  for each  $p \in P$ . We subtract  $\pi(v')$  from the NN distance to recover the reduced cost of the arc from  $v'$ .
3. Non-shortcut forward arcs from  $s$ -to- $A$  and from  $B$ -to- $t$ . For  $s$ , we can maintain a min-heap on the potentials of  $B \setminus \tilde{S}$ , queried only if  $v' = s$ . For  $B$ -to- $t$  arcs, there is only one arc to check if  $v' \in B$ , which we can examine manually.
4. Shortcut arcs  $(s, b)$  corresponding to null 2-paths from  $s$  to  $b \in (B \setminus B_\emptyset) \setminus S$ . For these, we maintain a NN data structure with  $P = A_\emptyset$ ,  $Q = (B \setminus B_\emptyset) \setminus S$  with weights  $\omega(p) = \pi(s)$  for all  $p \in P$ , and  $\omega(q) = \pi(q)$  for all  $q \in Q$ . A response  $(a, b)$  corresponds to th null 2-path  $(s, a, b)$ . This is only queried if  $v' = s$ .
5. Shortcut arcs  $(a, t)$  corresponding to null 2-paths from  $a \in (A \setminus A_\emptyset) \cap S$  to  $t$ . For these, we maintain a NN data structure over  $P = B_\emptyset \setminus S$  with weights  $\omega(p) = \pi(t)$  for each  $p \in P$ . A response  $(v', b)$  corresponds to th null 2-path  $(v', b, t)$ . We subtract  $\pi(v')$  from the NN distance to recover the reduced cost of the arc from  $v'$ . This is not queried if  $t \in \tilde{S}$ .
6. Shortcut arcs  $(s, t)$  corresponding to null 3-paths. For these, we maintain in a NN data structure with  $P = A_\emptyset \setminus S$ ,  $Q = B_\emptyset \setminus S$  with weights  $\omega(p) = \pi(s)$  for all  $p \in P$ , and  $\omega(q) = \pi(t)$  for all  $q \in Q$ . A response  $(a, b)$  corresponds to th null 3-path  $(s, a, b, t)$ . This is only queried while  $v' = s$  and  $t \notin S$ .

## C.2 Number of relaxations

First we bound the number of relaxations performed by both the Hungarian search and the depth-first search.

► **Lemma C.1.** *Hungarian search performs  $O(k)$  relaxations before a deficit vertex is reached.*

**Proof.** **«TO BE REWRITTEN.»** First we prove that there are  $O(k)$  non-shortcut relaxations. Each edge relaxation adds a new vertex to  $S$ , and non-shortcut relaxations only add normal vertices. The vertices of  $V \setminus S$  fall into several categories: (i)  $s$  or  $t$ , (ii) vertices of  $A$  or  $B$  with 0 imbalance, and (iii) deficit vertices of  $A$  or  $B$  ( $S$  contains all excess vertices). The number of vertices in (i) and (iii) is  $O(k)$ , leaving us to bound the number of (ii) vertices.

An  $A$  or  $B$  vertex with 0 imbalance must have an even number of  $\text{supp}(f)$  edges. There is either only one positive-capacity incoming arc (for  $A$ ) or outgoing arc (for  $B$ ), so this quantity is either 0 or 2. Since the vertex is normal, this must be 2. We charge 0.5 to each of the two

853  $\text{supp}(f)$  arcs; the arcs of  $\text{supp}(f)$  have no more than 1 charge each. Thus, the number of  
 854 type (ii) vertex relaxations is  $O(|\text{supp}(f)|)$ . By Corollary B.4,  $O(|\text{supp}(f)|) = O(k)$ .

855 Next we prove that there are  $O(k)$  shortcut relaxations. Recall the categories of shortcuts  
 856 from the list of data structures above. We have shortcuts corresponding to (i) null 2-paths  
 857 surrounding  $a \in A_\emptyset$ , (ii) null 2-paths surrounding  $b \in B_\emptyset$ , and (iii) null 3-paths, which go  
 858 from  $s$  to  $t$ . There is only one relaxation of type (iii), since  $t$  can only be added to  $S$  once.  
 859 The same argument holds for type (ii).

860 Each type (i) relaxation adds some normal  $b \in B \setminus B_\emptyset$  into  $S$ . Since  $b$  is normal, it must  
 861 either have deficit or an adjacent arc of  $\text{supp}(f)$ . We charge this relaxation to  $b$  if it is deficit,  
 862 or the adjacent arc of  $\text{supp}(f)$  otherwise. No vertex is charged more than once, and no  
 863  $\text{supp}(f)$  edge is charged more than twice, therefore the total number of type (i) relaxations  
 864 is  $O(|\text{supp}(f)|)$ . By Corollary B.4,  $O(|\text{supp}(f)|) = O(k)$ . ◀

865 Similarly we can prove that there are  $O(k)$  relaxations during the DFS.

866 ► **Corollary C.2.** *Depth-first search performs  $O(k)$  relaxations before a deficit vertex is*  
 867 *reached.*

### 868 C.3 Time analysis

869 Now we complete the time analysis by showing that each Hungarian search and depth-  
 870 first search can be implemented in  $O(k \text{ polylog } n)$  time after a one-time  $O(n \text{ polylog } n)$ -time  
 871 preprocessing.

872 There are only  $O(1)$  data structures for both. The rewinding mechanism and potential  
 873 update scheme of Section 2 to reset for each Hungarian search in time  $O(\text{polylog } n)$  times  
 874 the number of relaxations, giving the following. The main difference is that we are using  
 875 blocking flows rather than augmenting paths, so the initial  $S$  (i.e. excess vertices) may change  
 876 by more than one element for each REFINES iteration. The total number of such changes  
 877 throughout REFINES is  $O(k)$ , however, since REFINES begins with  $O(k)$  excess and does not  
 878 gain more. We obtain the following:

879 ► **Lemma C.3.** *After  $O(n \text{ polylog } n)$ -time preprocessing, each Hungarian search can be*  
 880 *implemented in  $O(k \text{ polylog } n)$  time.*

881 ► **Lemma C.4.** *After  $O(n \text{ polylog } n)$ -time preprocessing, each depth-first search can be*  
 882 *implemented in  $O(k \text{ polylog } n)$  time.*

883 ► **Lemma C.5.** *The rewinding mechanism takes in total  $O(k \text{ polylog } n)$  time in REFINES.*

### 884 C.4 Number of potential updates on null vertices

885 In our implementation of REFINES, we do not explicitly construct  $\tilde{H}_f$ ; instead we query its  
 886 edges using BCP/NN oracles and min/max heaps on elements of  $H_f$ . Potentials on the null  
 887 vertices are only required right before an augmentation sends a flow through a null path,  
 888 making the null vertices it passes normal. We use the construction from Lemma 4.1 to obtain  
 889 potential  $\pi$  on  $H_f$  such that the flow  $f$  is both  $\theta$ -optimal and admissible with respect to  $\pi$ .

890 **Size of blocking flows.** Now we bound the total number arcs whose flow is updated by a  
 891 blocking flow during the course of REFINES. This bounds both the time spent updating the  
 892 flow on these arcs and also the time spent on null vertex potential updates (Lemma C.7).

893 ► **Lemma C.6.** *The support of each blocking flow found in REFINES is of size  $O(k)$ .*

**Proof.** Let  $i$  be fixed and consider the invocation of DFS which produces the  $i$ -th blocking flow  $f_i$ . DFS constructs  $f_i$  as a sequence of admissible excess-deficit paths. Let one of these paths be  $P$ . Every arc in  $P$  is an arc relaxed by DFS, so  $N_i$  is bounded by the number of relaxations performed in DFS. Using Corollary C.2, we have  $N_i = O(k)$ . ◀

► **Lemma C.7.** *The number of end-of-REFINE null vertex potential updates is  $O(n)$ . The number of augmentation-induced null vertex potential updates in each invocation of REFINE is  $O(k \log k)$ .*

**Proof.** The number of end-of-REFINE potential updates is  $O(n)$ . Each update due to flow augmentation involves a blocking flow sending positive flow through a null path, causing a potential update on the passed null vertex. We charge this potential update to the edges of that null path, which are in turn arcs with positive flow in the blocking flow. For each blocking flow, no positive arc is charged more than twice. It follows that the number of augmentation-induced updates is at most the size of support of the blocking flow, which is  $O(k)$  by Lemma C.6. According to Lemma 3.4 there are  $O(\sqrt{k})$  iterations of REFINE before it terminates. Summing up we have an  $O(k\sqrt{k})$  bound over the course of REFINE. ◀

Now combining Lemma 3.3, Lemma 3.4, Lemma C.3, Lemma C.4, and Lemma C.7 completes the proof of Theorem 1.2.

## D Missing Details and Proofs from Section 5

### D.1 Uncapacitated MCF by excess scaling

We give an outline of the strongly polynomial-time algorithm for uncapacitated min-cost flow problem from Orlin [20]. Orlin's algorithm follows an *excess-scaling* paradigm originally due to Edmonds and Karp [10]. Consider the basic primal-dual framework used in the previous sections: The algorithm begins with both flow  $f$  and potentials  $\pi$  set to zero. Repeatedly run a *Hungarian search* that raises potentials (while maintaining dual feasibility) to create an admissible augmenting excess-deficit path, on which we perform flow augmentations. In terms of cost,  $f$  is maintained to be 0-optimal with respect to  $\pi$  and each augmentation over admissible edges preserves 0-optimality (by Lemma B.1). Thus, the final circulation must be optimal. The excess-scaling paradigm builds on top of this skeleton by specifying (i) between which excess and deficit vertices we send flows, and (ii) how much flow is sent by the augmentation.

The excess-scaling algorithm maintains a *scale parameter*  $\Delta$ , initially set to  $U$ . A vertex  $v$  with  $|\phi_f(v)| \geq \Delta$  is called *active*. Each augmenting path is chosen between an active excess vertex and an active deficit vertex. Once there are no more active excess or deficit vertices,  $\Delta$  is halved. Each sequence of augmentations where  $\Delta$  holds a constant value is called an *excess scale*. There are  $O(\log U)$  excess scales before  $\Delta < 1$  and, by integrality of supplies/demands,  $f$  is a circulation.

With some modifications to the excess-scaling algorithm, Orlin [20] obtains a strongly polynomial bound on the number of augmentations and excess scales. First, an *active* vertex is redefined to be one satisfying  $|\phi_f(v)| \geq \alpha\Delta$ , for a fixed parameter  $\alpha \in (0.5, 1)$ . Second, arcs with flow value at least  $3n\Delta$  at the beginning of a scale are *contracted* to create a new vertex, whose supply-demand is the sum of those on the two endpoints of the contracted arc. We use  $\hat{G} = (\hat{V}, \hat{E})$  to denote the resulting *contracted graph*, where each  $\hat{v} \in \hat{V}$  is a contracted component of vertices from  $V$ . Intuitively, the flow is so high on contracted arcs that no set of future augmentations can remove the arc from  $\text{supp}(f)$ . Third, in additional

to halving,  $\Delta$  is aggressively lowered to  $\max_{v \in V} \phi_f(v)$  if there are no active excess vertices and  $f(v \rightarrow w) = 0$  holds for every arc  $v \rightarrow w \in \hat{E}$ . Finally, flow values are not tracked within contracted components, but once an optimal circulation is found on  $\hat{G}$ , optimal potentials  $\pi^*$  can be *recovered* for  $G$  by sequentially undoing the contractions. The algorithm then performs a post-processing step which finds the optimal circulation  $f^*$  on  $G$  by solving a max-flow problem on the set of admissible arcs under  $\pi^*$ .

## D.2 Implementing contractions

«**REWRITE**» Following Agarwal *et al.* [2], our geometric data structures deals with real points in the plane instead of the contracted components. We will track the contracted components described in  $\hat{G}$  (e.g. with a disjoint-set data structure) and mark the arcs of  $\text{supp}(f)$  that are contracted. We maintain potentials on the points  $A$  and  $B$  directly, instead of the contracted components.

When conducting the Hungarian search, we initialize  $S$  to be the set of vertices from *active excess contracted components* who (in sum) meet the imbalance criteria. «**unclear**» Upon relaxing any  $v \in \hat{v}$ , we immediately relax all the contracted support arcs which span  $\hat{v}$ . Since the input network is uncapacitated, each contracted component is strongly connected in the residual network by the admissible forward/backward arcs of each contracted arc. «**unparsable**» To relax arcs in  $\hat{E}$ , we relax the support arcs before attempting to relax any non-support arcs; this will guarantee that the underlying graph of the support is acyclic (see Lemma D.6). Relaxations of support arcs can be performed without further potential changes, since they are admissible by invariant.

During the augmentations, contracted residual arcs are considered to have infinite capacity, and we do not update the value of flows on these arcs. We allow augmenting paths to begin from any point  $a \in \hat{v} \cap A$  in an active excess component  $\hat{v}$ , and end at any point  $b \in \hat{w} \cap B$  in an active deficit component  $\hat{w}$ .

## D.3 Recovering the optimal flow

We use the recovery strategy from Agarwal *et al.* [2], which runs in  $O(n \text{polylog } n)$  time. The main idea is that, if  $\mathcal{T}$  is an undirected *spanning tree of admissible edges* under optimal potentials  $\pi^*$ , then there exists an optimal flow  $f^*$  with support only on arcs corresponding to edges of  $\mathcal{T}$ . Intuitively,  $\mathcal{T}$  is a maximal set of linearly independent dual LP constraints for the optimal dual ( $\pi^*$ ), so there exists an optimal primal solution ( $f^*$ ) with support only on these arcs. To see this, we can use a perturbation argument: raising the cost of each non-tree edge by a positive amount does not change  $\text{cost}(\pi^*)$  or the feasibility of  $\pi^*$ , but does raise the cost of any circulation  $f$  using non-tree edges. Strong duality suggests that  $\text{cost}(f^*) = \text{cost}(\pi^*)$  is unchanged, therefore  $f^*$  must have support only on the tree edges.

Since the arcs corresponding to edges of  $\mathcal{T}$  have no cycles, we can solve the maximum flow in linear time using the following greedy algorithm. Let  $\text{par}(v)$  be the parent of vertex  $v$  in  $\mathcal{T}$ . We begin with  $f^* = 0$  and process  $\mathcal{T}$  from its leaves upwards. For a supply leaf  $v$ , we satisfy its supply by choosing  $f^*(v \rightarrow \text{par}(v)) \leftarrow \phi(v)$ . Otherwise if leaf  $v$  is a demand vertex, we choose  $f^*(\text{par}(v) \rightarrow v) \leftarrow -\phi(v)$ . Once we've solved the supplies/demands for each leaf, then we can *trim* the leaves, removing them from  $\mathcal{T}$  and setting the supply/demand of each parent-of-a-leaf to its current imbalance. Then, we can recurse on this smaller tree and its new set of leaves.

► **Lemma D.1.** *Let  $G(\mathcal{T})$  be the subnetwork of  $G$  corresponding to edges of the undirected spanning tree  $\mathcal{T}$ . If there exists a flow in  $G(\mathcal{T})$  which satisfies every supply and demand, then*

983 the greedy algorithm finds the maximum flow in  $G(\mathcal{T})$  in  $O(n)$  time.

984 **Proof.** Observe that, for any flow  $f$  in  $G$ ,  $\text{supp}(f)$  has no paths of length longer than one.  
 985 Thus, if a flow  $f^*$  satisfying supplies/demands exists within  $G(\mathcal{T})$ , then each supply vertex has  
 986 flow paths that terminate at its parent/children. Similarly, each demand vertex receive all its  
 987 flow from its parent/children. Since there is only one option for a supply leaf (resp. demand  
 988 leaf) to send its flow (resp. receive its flow), the greedy algorithm correctly identifies the  
 989 values of  $f^*$  for arcs adjoining  $\mathcal{T}$  leaves. Trimming these leaves, we can apply this argument  
 990 recursively for their parents. The running time of the greedy algorithm is  $O(n)$ , as leaves  
 991 can be identified in  $O(n)$  time and no vertex becomes a leaf more than once. ◀

992 It remains to show how we construct  $\mathcal{T}$ . We begin with a (spanning) *shortest path tree*  
 993 (SPT)  $T$  in the residual network of  $f$ , under reduced costs and rooted at an arbitrary vertex  
 994  $r$ . For the SPT to span, we need the additional assumption that  $G$  is strongly connected.  
 995 We make can make  $G$  strongly connected by adding a 0-supply vertex  $s$  with arcs  $s \rightarrow a$  for all  
 996  $a \in A$  and  $b \rightarrow a$  for all  $b \in B$ , with some high cost  $M$ . Following Orlin [20], these arcs cannot  
 997 appear in an optimal flow if  $M$  is sufficiently high, and we can extend  $\pi^*$  to include  $s$  using  
 998  $\pi^*(s) = 0$  if  $M > \max_{b \in B} \pi^*(b)$ . This extension to  $\pi^*$  preserves feasibility.

999 The edges corresponding to arcs of  $T$  do not suffice for  $\mathcal{T}$ , since some SPT arcs may be  
 1000 inadmissible. Let  $d_r(v)$  be the shortest path distance of  $v \in A \cup B \cup \{s\}$  from  $r$ , and consider  
 1001 potentials  $\pi^\# = \pi^* - d_r$ .

1002 ▶ **Lemma D.2 (Orlin [20, Lemma 3]).** *Let  $f$  be a flow satisfying the optimality conditions*  
 1003 *with respect to  $\pi^*$ . Then, (i)  $f$  satisfies the optimality conditions with respect to  $\pi^\#$ , and (ii)*  
 1004 *all SPT arcs are admissible under  $\pi^\#$ .*

1005 We can use this lemma to argue that  $\pi^\#$  is still optimal. Recall that  $f$  has values defined  
 1006 only on the non-contracted residual arcs; we can apply the first part of Lemma D.2 on these  
 1007 arcs. For arcs within contracted components, we use a different argument. Observe that  
 1008 each  $\hat{v} \in \hat{V}$  is spanned by a set of  $\text{supp}(f)$  arcs, which are admissible by invariant. Thus, all  
 1009  $v \in \hat{v}$  are equidistant from  $r$ , and they will have the same value  $d_r(v)$ . It follows that the  
 1010 reduced costs of arcs with both endpoints in  $\hat{v}$  do not change when replacing  $\pi^*$  with  $\pi^\#$ , so  
 1011 arcs contained in  $\hat{v}$  that met the optimality conditions for  $\pi^*$  still meet them for  $\pi^\#$ .

1012 From the second part of Lemma D.2, the SPT  $T$  is a spanning tree of admissible arcs  
 1013 under  $\pi^\#$ . We set  $\mathcal{T}$  to be the set of undirected edges corresponding to  $T$ .

1014 **Computing the SPT.** We conclude by describing the procedure for building the SPT, i.e.  
 1015 by running Dijkstra's algorithm in the residual network. We use a geometric implementation  
 1016 that is very similar to Hungarian search. We begin with  $S = \{r\}$  and  $d_r(r) = 0$ , where  $r$  is  
 1017 our arbitrary root. For all other vertices,  $d_r(v)$  is initially unknown. In each iteration, we  
 1018 relax the minimum-reduced cost arc  $v \rightarrow w$  in the frontier  $S \times (A \cup B) \setminus S$ , adding  $w$  to  $S$ , and  
 1019 setting  $d_r(w) = d_r(v) + c_{\pi^*}(v, w)$ . Once  $S = A \cup B$ , the SPT  $T$  is the set of relaxed arcs.

1020 If an either direction of an arc of  $\text{supp}(f)$  enters the frontier, we relax it immediately.  
 1021 To detect support arcs, we build a list for each  $v \in A \cup B$  of the support arcs which use  $v$   
 1022 as an endpoint, and once  $v \in S$  we check its list. There are  $O(n)$  support arcs in total (by  
 1023 acyclicity of  $E(\text{supp}(f))$ ; see Lemma D.6), so the total time spent searching these lists is  
 1024  $O(n)$ . Such relaxations are correct for the shortest path tree, since the support edges are  
 1025 admissible and reduced costs are nonnegative.

1026 Other edges appearing in the frontier can be split into three categories:

1027 1. Forward  $A$ -to- $B$  arcs. We query these using a BCP with  $P = A \cap S$  and  $Q = B \setminus S$ .

1028 2.  $B$ -to- $s$  arcs. These will never have flow support. We can query the minimum with a  
 1029 max-heap on potentials of  $B \cap S$ . We query these while  $s \in S$ .  
 1030 3.  $s$ -to- $A$  arcs. These will also never have flow support. We can query the minimum with a  
 1031 min-heap on potentials of  $A \setminus S$ . We query these while  $s \in S$ .  
 1032 We perform  $O(n)$  relaxations and takes  $O(\text{polylog } n)$  time per relaxation, for non-support  
 1033 relaxations. An additional  $O(n)$  time is spent relaxing support edges. The total running  
 1034 time of Dijkstra's algorithm is  $O(n \text{ polylog } n)$ . Combining with Lemma D.1, we obtain the  
 1035 following.

1036 ► **Lemma D.3.** *Given optimal potentials  $\pi^*$  and an optimal contracted flow  $f$ , the optimal*  
 1037 *flow  $f^*$  can be computed in  $O(n \text{ polylog } n)$  time.*

#### 1038 D.4 Recovering the optimal flow for sum-of-distances.

1039 When the matching objective uses the just the  $p$ -norms (that is, when  $q = 1$ ), we can prove  
 1040 that the subgraph formed by admissible arcs is in fact *planar*. Planarity gives us two things  
 1041 for recovery: there are only a linear number of admissible arcs, and the max-flow on them  
 1042 can be solved in near-linear time with a planar graph multiple-source multiple-sink max-flow  
 1043 algorithm. Note that this recovery algorithm is not asymptotically faster than the other, and  
 1044 depending on the planar max-flow algorithm, not necessarily simpler either.

1045 Up until now, we have not placed restrictions on coincidence between  $A$  and  $B$ , but  
 1046 for the next proof it is useful to do so. We can assume that all points within  $A \cup B$  are  
 1047 distinct, otherwise we can replace all points coincident at  $x \in \mathbb{R}^2$  with a single point whose  
 1048 supply/demand is  $\sum_{v \in A \cup B: v=x} \lambda(v)$ . This is roughly equivalent to transporting as much as  
 1049 we can between coincident supply and demand, and is optimal by triangle inequality.

1050 Without loss of generality, assume  $\pi^*$  is nonnegative (raising  $\pi^*$  uniformly on all points  
 1051 does not change the objective or feasibility). Recall that  $\pi^*$  is feasibility if for all  $a \in A$  and  
 1052  $b \in B$

$$1053 \quad c_{\pi^*}(a \rightarrow b) = \|a - b\|_p - \pi^*(a) + \pi^*(b) \geq 0.$$

1054 An arc  $a \rightarrow b$  is admissible when

$$1055 \quad c_{\pi^*}(a \rightarrow b) = \|a - b\|_p - \pi^*(a) + \pi^*(b) = 0.$$

1056 We note that these definitions have a nice visual: Place disks  $D_q$  of radius  $\pi(q)$  at each  
 1057  $q \in A \cup B$ . Feasibility states that for all  $a \in A$  and  $b \in B$ ,  $D_a$  cannot contain  $D_b$  with a  
 1058 gap between their boundaries. The arc  $a \rightarrow b$  is admissible when  $D_a$  contains  $D_b$  and their  
 1059 boundaries are tangent.

1060 ► **Lemma D.4.** *Let  $\pi^*$  be a set of optimal potentials for the point sets  $A$  and  $B$ , under costs*  
 1061  *$c(a, b) = \|a - b\|_p$ . Then, the set of admissible arcs under  $\pi^*$  form a planar graph.*

1062 **Proof.** We assume the points of  $A \cup B$  are in general position (e.g. by symbolic perturbation)  
 1063 such that no three points are collinear. Let  $a_1 \rightarrow b_1$  and  $a_2 \rightarrow b_2$  be any pair of admissible  
 1064 arcs under  $\pi^*$ . We will isolate them from the rest of the points, considering  $\pi^*$  restricted  
 1065 to the four points  $\{a_1, a_2, b_1, b_2\}$ . Clearly, this does not change whether the two arcs cross.  
 1066 Observe that we can raise  $\pi^*(a_2)$  and  $\pi^*(b_2)$  uniformly, until  $c_{\pi^*}(a_2, b_1) = 0$ , without breaking  
 1067 feasibility or changing admissibility of  $a_1 \rightarrow b_1$  and  $a_2 \rightarrow b_2$ . Henceforth, we assume that we have  
 1068 modified  $\pi^*$  in this way to make  $a_2 \rightarrow b_1$  admissible. Given positions of  $a_1$ ,  $a_2$ , and  $b_1$ , we now  
 1069 try to place  $b_2$  such that  $a_1 \rightarrow b_1$  and  $a_2 \rightarrow b_2$  cross. Specifically,  $b_2$  must be placed within a  
 1070 region  $\mathcal{F}$  that lies between the rays  $\overrightarrow{a_2 a_1}$  and  $\overrightarrow{a_2 b_1}$ , and within the halfplane bounded by  $\overleftrightarrow{a_1 b_1}$   
 1071 that does not contain  $a_2$ .



Let  $g_a(q) := \|a - q\| - \pi^*(a)$  for  $a \in A$  and  $q \in \mathbb{R}^2$ . Let the *bisector* between  $a_1$  and  $a_2$  be  $\beta := \{q \in \mathbb{R}^2 \mid g_{a_1}(q) = g_{a_2}(q)\}$ .  $\beta$  is a curve subdividing the plane into two open faces, one where  $g_{a_1}$  is minimized and the other where  $g_{a_2}$  is. From these definitions, admissibility of  $a_1 \rightarrow b_1$  and  $a_2 \rightarrow b_1$  imply that  $b_1$  is a point of the bisector.

We show that  $\mathcal{F}$  lies entirely on the  $g_{a_1}$  side of the bisector. First, we prove that the closed segment  $\overline{a_1 b_1}$  lies entirely on the  $g_{a_1}$  side, except  $b_1$  which lies on  $\beta$ . Any  $q \in \overline{a_1 b_1}$  can be written parametrically as  $q(t) = (1 - t)b_1 + ta_1$  for  $t \in [0, 1]$ . Consider the single-variable functions  $g_{a_1}(q(t))$  and  $g_{a_2}(q(t))$ .

$$\begin{aligned} g_{a_1}(q(t)) &= (1 - t)\|a_1 - b_1\| - \pi(a_1) \\ g_{a_2}(q(t)) &= \|(a_2 - b_1) - t(a_1 - b_1)\| - \pi(a_2) \end{aligned}$$

At  $t = 0$ , these expressions are equal. Observe that the derivative with respect to  $t$  of  $g_{a_1}(q(t))$  is less than  $g_{a_2}(q(t))$ . Indeed, the value of  $\frac{d}{dt}\|(a_2 - b_1) - t(a_1 - b_1)\|$  is at least  $-\|a_1 - b_1\| = \frac{d}{dt}g_{a_1}(q(t))$ , which is realized if and only if  $\frac{(a_2 - b_1)}{\|a_2 - b_1\|} = \frac{(a_1 - b_1)}{\|a_1 - b_1\|}$ . This corresponds to  $\overrightarrow{a_2 b_1}$  and  $\overrightarrow{a_1 b_1}$  being parallel, but this is disallowed since  $a_1, a_2, b_1$  are in general position. Thus,  $g_{a_1}(q(t)) \leq g_{a_2}(q(t))$  with equality only at  $b_1$ .

Now, we parameterize each point of  $\mathcal{F}$  in terms of points on  $\overline{a_1 b_1}$ . Every  $q \in \mathcal{F}$  can be written as  $q(t') = q' + t'(q' - a_2)$  for some  $q' \in \overline{a_1 b_1}$  and  $t \geq 0$ , i.e.  $q' = \overline{a_1 b_1} \cap \overrightarrow{a_2 q}$ . We call  $q'$  the *projection* of  $q$  onto  $\overline{a_1 b_1}$ . We can write  $g_{a_1}$  and  $g_{a_2}$  in terms of  $t'$  and observe that  $\frac{d}{dt'}g_{a_1}(q(t')) \leq \frac{d}{dt'}g_{a_2}(q(t'))$ , as the derivative of  $g_a(q(t'))$  is maximized if  $(q(t') - a)$  is parallel to  $(q(t') - a_2)$  and lower otherwise. Notably,  $q(t')$  with projection  $b_1$  have  $\frac{d}{dt'}g_{a_1}(q(t')) < \frac{d}{dt'}g_{a_2}(q(t'))$ , since  $a_1, a_2, b_1$  are in general position. Any  $q(t')$  with a different projection do not have strict inequality, but the projection itself has  $g_{a_1}(q') < g_{a_2}(q')$  for  $q' \neq b_1$  since it lies on  $\overline{a_1 b_1}$ . Therefore, for all  $q \in \mathcal{F} \setminus \{b_1\}$ ,  $g_{a_1}(q') < g_{a_2}(q')$ , and  $\mathcal{F}$  lies on the  $g_{a_1}$  side of the bisector except for  $b_1$  which lies on  $\beta$ . We can eliminate  $b_1$  as a candidate position for  $b_2$ , since points of  $B$  cannot coincide.

Observe that  $g_{a_1}(b) < g_{a_2}(b)$  for  $b \in B$  implies that  $c_\pi(a_1, b) < c_\pi(a_2, b)$ , and  $c_\pi(a_1, b) = c_\pi(a_2, b)$  if and only if  $b$  lies on  $\beta$ . This holds for all  $b \in \mathcal{F}$  including our prospective  $b_2$ , but then  $c_\pi(a_1, b_2) < c_\pi(a_2, b_2) = 0$  since  $a_2 \rightarrow b_2$  is admissible. This violates feasibility of  $a_1 \rightarrow b_2$ , so there is no feasible placement of  $b_2$  which also crosses  $a_1 \rightarrow b_1$  with  $a_2 \rightarrow b_2$ . ◀

We can construct the entire set of admissible arcs by repeatedly querying the minimum-reduced-cost outgoing arc for each  $a \in A$  until the result is not admissible. By Lemma D.4 the resulting arc set forms a planar graph, so by Euler's formula the number of arcs to query is  $O(n)$ . We can then find the maximum flow in time  $O(n \log^3 n)$  time, using the planar multiple-source multiple-sink maximum-flow algorithm by Borradaile *et al.* [7].

► **Lemma D.5.** *If the transportation objective is sum-of-costs, then given the optimal potentials  $\pi^*$ , we can compute an optimal flow  $f^*$  in  $O(n \text{ polylog } n)$  time.*

## D.5 Dead vertices

Let the *support degree* of a vertex be its degree in the graph induced by the underlying edges of  $\text{supp}(f)$ . We call a vertex  $b \in B$  *dead* if  $b$  has support degree 0 and is not an active excess or deficit vertex; call it *living* otherwise. Dead vertices are essentially equivalent to the *null vertices* of Section 3. However, since the reduction in this section does not use a super-source/super-sink, we can simply remove these from consideration during a Hungarian search — they will not terminate the search, and have no outgoing residual arcs. Like the null vertices, we ignore dead vertex potentials and infer feasible potentials when they become

live again. We use  $A_\ell$  and  $B_\ell$  to denote the living vertices of points in  $A$  and  $B$ , respectively. Note that being dead/alive is a notion strictly defined only for vertices, and not for contracted components.

We say a dead vertex is *revived* when it stops meeting either condition of the definition. Dead vertices are only revived after  $\Delta$  decreases (at the start of a subsequent excess scale) as no augmenting path will cross a dead vertex and they cannot meet the criteria for contractions. When a dead vertex is revived, we must add it back into each of our data structures and give it a feasible potential. For revived  $b \in B$ , a feasible choice of potential is  $\pi(b) \leftarrow \max_{a \in A} (\pi(a) - c(a, b))$  which we can query by maintaining a weighted nearest neighbor data structure on the points of  $A$ . The total number of revivals is bounded above by the number of augmentations: since the final flow is a circulation on  $\hat{G}$  and a newly revived vertex  $v$  has no incident arcs in  $\text{supp}(f)$  and cannot be contracted, there is at least one subsequent augmentation which uses  $v$  as its beginning or end. Thus, the total number of revivals is  $O(n \log n)$ .

## D.6 Number of relaxations

By prioritizing the relaxation of support arcs, we also have the following lemma.

► **Lemma D.6 (Agarwal et al. [2]).** *If arcs of  $\text{supp}(f)$  are relaxed first as they arrive on the frontier, then  $E(\text{supp}(f))$  is acyclic.*

**Proof.** Let  $f_i$  be the pseudoflow after the  $i$ -th augmentation, and let  $T_i$  be the forest of relaxed arcs generated by the Hungarian search for the  $i$ -th augmentation. Namely, the  $i$ -th augmenting path is an excess-deficit path in  $T_i$ , and all arcs of  $T_i$  are admissible by the time the augmentation is performed. Let  $E(T_i)$  be the undirected edges corresponding to arcs of  $T_i$ . Notice that,  $E(\text{supp}(f_{i+1})) \subseteq E(\text{supp}(f_i)) \cup E(T_i)$ . We prove that  $E(\text{supp}(f_i)) \cup E(T_i)$  is acyclic by induction on  $i$ ; as  $E(\text{supp}(f_{i+1}))$  is a subset of these edges, it must also be acyclic. At the beginning with  $f_0 = 0$ ,  $E(\text{supp}(f_0))$  is vacuously acyclic.

Let  $E(\text{supp}(f_i))$  be acyclic by induction hypothesis. Since  $T_i$  is a forest (thus, acyclic), any hypothetical cycle  $\Gamma$  that forms in  $E(\text{supp}(f_i)) \cup E(T_i)$  must contain edges from both  $E(\text{supp}(f_i))$  and  $E(T_i)$ . To give a visual analogy, we will color  $e \in \Gamma$  *purple* if  $e \in E(\text{supp}(f_i)) \cap E(T_i)$ , *red* if  $e \in E(\text{supp}(f_i))$  but  $e \notin E(T_i)$ , and *blue* if  $e \in E(T_i)$  but  $e \notin E(\text{supp}(f_i))$ . Then,  $\Gamma$  is neither entirely red nor entirely blue. We say that red and purple edges are *red-tinted*, and similarly blue and purple edges are *blue-tinted*. Roughly speaking, our implementation of the Hungarian search prioritizes relaxing red-tinted admissible arcs over pure blue arcs.

We can sort the blue-tinted edges of  $\Gamma$  by the order they were relaxed into  $S$  during the Hungarian search forming  $T_i$ . Let  $(v, w) \in \Gamma$  be the last pure blue edge relaxed, of all the blue-tinted edges in  $\Gamma$  — after  $(v, w)$  is relaxed, the remaining unrelaxed, blue-tinted edges of  $\Gamma$  are purple.

Let us pause the Hungarian search the moment before  $(v, w)$  is relaxed. At this point,  $v \in S$  and  $w \notin S$ , and the Hungarian search must have finished relaxing all frontier support arcs. By our choice of  $(v, w)$ ,  $\Gamma \setminus (v, w)$  is a path of relaxed blue edges and red-tinted edges which connect  $v$  and  $w$ . Walking around  $\Gamma \setminus (v, w)$  from  $v$  to  $w$ , we see that every vertex of the cycle must be in  $S$  already:  $v \in S$ , relaxed blue edges have both endpoints in  $S$ , and any unrelaxed red-tinted edge must have both endpoints in  $S$ , since the Hungarian search would have prioritized relaxing the red-tinted edges to grow  $S$  before relaxing  $(v, w)$  (a blue edge). It follows that  $w \in S$  already, a contradiction.

1160 No such cycle  $\Gamma$  can exist, thus  $E(\text{supp}(f_i)) \cup E(T_i)$  is acyclic and  $E(\text{supp}(f_{i+1})) \subseteq$   
 1161  $E(\text{supp}(f_i)) \cup E(T_i)$  is acyclic. By induction,  $E(\text{supp}(f_i))$  is acyclic for all  $i$ . ◀

1162 Let  $E(\Sigma_a)$  **◀only used once▶** be the underlying edges of the support star centered at  $a$   
 1163 and  $F := E(\text{supp}(f)) \setminus \bigcup_{a \in A} E(\Sigma_a)$ . Using Lemma D.6, we can show that the number of  
 1164 support arcs outside support stars ( $|F|$ ) is small.

1165 ▶ **Lemma D.7.**  $|B_\ell \setminus \bigcup_{a \in A} \Sigma_a| \leq r$ .

1166 **Proof.**  $F$  is constructed from  $E(\text{supp}(f))$  by eliminating edges in support stars, therefore all  
 1167 edges in  $F$  must adjoin vertices in  $B$  of support degree at least 2. By Lemma D.6,  $E(\text{supp}(f))$   
 1168 is acyclic and therefore forms a spanning forest over  $A \cup B_\ell$ , so  $F$  is also a bipartite forest.  
 1169 All leaves of  $F$  are therefore vertices of  $A$ .

1170 Pick an arbitrary root for each connected component of  $F$  to establish parent-child  
 1171 relationships for each edge. As no vertex in  $B$  is a leaf, each vertex in  $B$  has at least one  
 1172 child. Charge each vertex in  $B$  to one of its children in  $F$ , which must belong to  $A$ . Each  
 1173 vertex in  $A$  is charged at most once. Thus, the number of  $B_\ell$  vertices outside of support  
 1174 stars is no more than  $r$ . ◀

1175 ▶ **Lemma 5.2.** *Suppose we have stripped the graph of dead vertices. The number of relaxation*  
 1176 *steps in a Hungarian search outside of support stars is  $O(r)$ .*

1177 **Proof.** If there are no dead vertices, then each non-support star relaxation step adds either  
 1178 (i) an active deficit vertex, (ii) a non-deficit vertex  $a \in A_\ell$ , or (iii) a non-deficit vertex  $b \in B_\ell$   
 1179 of support degree at least 2. There is a single relaxation of type (i), as it terminates the  
 1180 search. The number of vertices of type (ii) is  $r$ , and the number of vertices of type (iii) is at  
 1181 most  $r$  by Lemma D.7. The lemma follows. ◀

1182 ▶ **Lemma 5.3.** *Hungarian search takes  $O(r\sqrt{n} \text{polylog } n)$  time.*

1183 **Proof.** The number of relaxation steps outside of support stars is  $O(r)$  by Lemma 5.2. The  
 1184 time per relaxation outside of support stars is  $O(\sqrt{n} \text{polylog } n)$ . The time spent processing  
 1185 relaxations within a support star is  $O(\sqrt{n} \text{polylog } n)$ , and at most  $r$  are relaxed during the  
 1186 search. The total time is therefore  $O(r\sqrt{n} \text{polylog } n)$ . ◀

## 1187 D.7 Updating support stars

1188 Initially, we label stars big or small according to the  $\sqrt{n}$  threshold. Afterwards, a star that is  
 1189 currently big is turned into a small star once  $|\Sigma_a| \leq \sqrt{n}/2$ , and star that is currently small  
 1190 is turned into a big star once  $|\Sigma_a| \geq 2\sqrt{n}$ . We say a star which crosses one of these size  
 1191 thresholds is **changing state** (from small-to-big or big-to-small), and must be represented in  
 1192 the opposite type of data structure. Our strategy is to charge the data structure update time  
 1193 associated with a state change to the **membership changes** in  $\Sigma_a$  that preceded the state  
 1194 change.

1195 A star  $\Sigma_a$  undergoing a big-to-small state change has size  $|\Sigma_a| \leq \sqrt{n}/2$ . The state change  
 1196 deletes  $\mathcal{D}_{\text{big}}(a)$  and inserts  $\Sigma_a$  into  $\mathcal{D}_{\text{small}}$ . Thus, the time spent for a big-to-small state  
 1197 change is  $O(\sqrt{n} \text{polylog } n)$ , and there were at least  $\sqrt{n}/2$  points removed from  $\Sigma_a$  since it  
 1198 last changed state. The amortized time for a big-to-small state change per star membership  
 1199 change is  $O(\text{polylog } n)$ .

1200 A star  $\Sigma_a$  undergoing a small-to-big state change has size  $|\Sigma_a| \geq 2\sqrt{n}$ . We can write its  
 1201 size as  $|\Sigma_a| = \sqrt{n} + x$  for some integer  $x \geq \sqrt{n}$ , so we also have  $|\Sigma_a| \leq 2x$ . When switching,  
 1202 we delete all  $|\Sigma_a|$  points from  $\mathcal{D}_{\text{small}}$  and construct a new  $\mathcal{D}_{\text{big}}(a)$ . Constructing  $\mathcal{D}_{\text{big}}(a)$

1203 requires inserting up to  $r$  points of  $A$  (into  $P$ ) and the  $|\Sigma_a|$  points of the star (into  $Q$ ). Thus,  
 1204 the time spent for a small-to-big state change is  $(r + 2x) \cdot O(\text{polylog } n)$ , and there were at  
 1205 least  $x$  points added to  $\Sigma_a$  since it last changed state. The amortized time for a small-to-big  
 1206 state change per star membership change is  $O((r/x) \text{polylog } n)$ . Since  $x \geq \sqrt{n}$ , this is at  
 1207 most  $O((r/\sqrt{n}) \text{polylog } n)$ .

1208      Star membership can only be changed by augmenting paths passing through the vertex,  
 1209 therefore the total number of membership changes is  $O(rn \log n)$  by Lemmas 5.1 and 5.2.  
 1210 Thus, the total time spent on state changes is  $O((r^2 \sqrt{n} + rn) \text{polylog } n)$ .