

Efficient Algorithms for Geometric Partial Matching

Pankaj K. Agarwal

Duke University, USA

pankaj@cs.duke.edu

Hsien-Chih Chang

Duke University, USA

hsienchih.chang@duke.edu

Allen Xiao

Duke University, USA

axiao@cs.duke.edu

1 Abstract

Let A and B be two point sets in the plane of sizes r and n respectively (assume $r \leq n$), and let k be a parameter. A matching between A and B is a family of pairs in $A \times B$ so that any point of $A \cup B$ appears in at most one pair. Given two positive integers p and q , we define the cost of matching M to be $c(M) = \sum_{(a,b) \in M} \|a - b\|_p^q$ where $\|\cdot\|_p$ is the L_p -norm. The geometric partial matching problem asks to find the minimum-cost size- k matching between A and B .

We present efficient algorithms for geometric partial matching problem that work for any powers of L_p -norm matching objective: An exact algorithm that runs in $O((n + k^2) \text{polylog } n)$ time, and a $(1 + \varepsilon)$ -approximation algorithm that runs in $O((n + k\sqrt{k}) \text{polylog } n \cdot \log \varepsilon^{-1})$ time. Both algorithms are based on the primal-dual flow augmentation scheme; the main improvements are obtained by using dynamic data structures to achieve efficient flow augmentations. Using similar techniques, we give an exact algorithm for the planar transportation problem that runs in $O(\min\{n^2, rn^{3/2}\} \text{polylog } n)$ time.

2012 ACM Subject Classification Theory of computation → Design and analysis of algorithms

Keywords and phrases partial matching, transportation, optimal transport, minimum-cost flow, bichromatic closest pair

Funding Work on this paper was supported by NSF under grants CCF-15-13816, CCF-15-46392, and IIS-14-08846, by an ARO grant W911NF-15-1-0408, and by BSF Grant 2012/229 from the U.S.-Israel Binational Science Foundation.

Acknowledgements We thank Haim Kaplan for discussion and suggestion to use Goldberg *et al.* [8] algorithm. We thank Debmalya Panigrahi for helpful discussions.

Lines 533

14 1 Introduction

Given two point sets A and B in the plane, we consider the problem of finding the minimum-cost partial matching between A and B . Formally, suppose A has size r and B has size n where $r \leq n$. Let $G(A, B)$ be the undirected complete bipartite graph between A and B , and let the cost of edge (a, b) be $c(a, b) = \|a - b\|_p^q$, for some positive integers p and q . A matching M in $G(A, B)$ is a set of edges sharing no endpoints. The size of M is the number of edges in M . The cost of matching M , denoted $c(M)$, is defined to be the sum of costs of



© Pankaj K. Agarwal, Hsien-Chih Chang, Allen Xiao;
licensed under Creative Commons License CC-BY

The 35th International Symposium on Computational Geometry (SOCG 2019).



Leibniz International Proceedings in Informatics

Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany



edges in M . For a parameter k , the problem of finding the minimum-cost size- k matching in $G(A, B)$ is called the *geometric partial matching problem*. We call the corresponding problem in general bipartite graphs (with arbitrary edge costs) the *partial matching problem*.¹

We also consider the following generalization of bipartite matching. Let $\phi : A \cup B \rightarrow \mathbb{Z}$ be an integral *supply-demand function* with positive value on points of A and negative value on points of B , satisfying $\sum_{a \in A} \phi(a) = -\sum_{b \in B} \phi(b)$. A *transportation map* is a function $\tau : A \times B \rightarrow \mathbb{R}_{\geq 0}$ such that $\sum_{b \in B} \tau(a, b) = \phi(a)$ for all $a \in A$ and $\sum_{a \in A} \tau(a, b) = -\phi(b)$ for all $b \in B$. We define the cost of τ to be

$$c(\tau) := \sum_{(a,b) \in A \times B} c(a, b) \cdot \tau(a, b).$$

The *transportation problem* asks to compute a transportation map of minimum cost.

Related work. Maximum-size bipartite matching is a classical problem in graph algorithms. Upper bounds include the $O(m\sqrt{n})$ time algorithm by Hopcroft and Karp [9] and the $O(m \min\{\sqrt{m}, n^{2/3}\})$ time algorithm by Even and Tarjan [6], where n is the number of vertices and m is the number of edges. The first improvement in over thirty years was made by Mądry [12], which uses an interior-point algorithm, runs in $O(m^{10/7} \text{polylog } n)$ time.

The Hungarian algorithm [11] computes a minimum-cost maximum matching in a bipartite graph in roughly $O(mn)$ time. Faster algorithms have been developed, such as the $O(m\sqrt{n} \log(nC))$ time algorithms by Gabow and Tarjan [7] and the improved $O(m\sqrt{n} \log C)$ time algorithm by Duan *et al.* [5] assuming the edge costs are integral; here C is the maximum cost of an edge. Ramshaw and Tarjan [14] showed that the Hungarian algorithm can be extended to compute a minimum-cost partial matching of size k in $O(km + k^2 \log r)$ time, where r is the size of the smaller side of the bipartite graph. They also proposed a cost-scaling algorithm for partial matching that runs in time $O(m\sqrt{k} \log(kC))$, again assuming that costs are integral. By reduction to unit-capacity min-cost flow, Goldberg *et al.* [8] developed a cost-scaling algorithm for partial matching with an identical running time $O(m\sqrt{k} \log(kC))$, again only for integral edge costs.

In geometric settings, the Hungarian algorithm can be implemented to compute an optimal perfect matching between A and B (assuming equal size) in time $O(n^2 \text{polylog } n)$ [10] (see also [1, 18]). This algorithm computes an optimal size- k matching in time $O(kn \text{polylog } n)$. Faster approximation algorithms have been developed for computing perfect matchings in geometric settings [3, 15, 18, 19]. Recall that the cost of the edges are the q th power of their L_p -distances. When $q = 1$, the best algorithm to date by Sharathkumar and Agarwal [16] computes a $(1 + \varepsilon)$ -approximation to the optimal perfect matching in $O(n \text{polylog } n \cdot \text{poly } \varepsilon^{-1})$ expected time with high probability. Their algorithm can also compute a $(1 + \varepsilon)$ -approximate partial matching within the same time bound. For $q > 1$, the best known approximation algorithm to compute a perfect matching runs in $O(n^{3/2} \text{polylog } n \log(1/\varepsilon))$ time [15]; it is not obvious how to extend this algorithm to the partial matching setting.

The transportation problem can also be formulated as an instance of the minimum-cost flow problem. The strongly polynomial uncapacitated min-cost flow algorithm by Orlin [13] solves the transportation problem in $O((m + n \log n)n \log n)$ time. Lee and Sidford [?] give a weakly polynomial algorithm that runs in $O(m\sqrt{n} \text{polylog}(n, U))$ time, where U is the maximum amount of vertex supply-demand. Agarwal *et al.* [2] showed that Orlin's algorithm can be implemented to solve 2D transportation in time $O(n^2 \text{polylog } n)$. By adapting the

¹ Partial matching is also called *imperfect matching* or *imperfect assignment* [8, 14].

65 Lee-Sidford algorithm, they developed a $(1 + \varepsilon)$ -approximation algorithm that runs in
 66 $O(n^{3/2}\varepsilon^{-2} \text{polylog}(n, U))$ time. They also gave a Monte-Carlo algorithm that computes
 67 an $O(\log^2(1/\varepsilon))$ -approximate solution in $O(n^{1+\varepsilon})$ time with high probability. **«cite the**
 68 **preconditioning paper?»**

69 **Our results.** There are three main results in this paper. First in Section 2 we present an
 70 efficient algorithm for computing an optimal partial matching in \mathbb{R}^2 .

71 **► Theorem 1.1.** *Given two point sets A and B in \mathbb{R}^2 each of size at most n and an*
 72 *integer $k \leq n$, a minimum-cost matching of size k between A and B can be computed in*
 73 *$O((n + k^2) \text{polylog } n)$ time.*

74 We use *bichromatic closest pair (BCP)* data structures to implement the Hungarian
 75 algorithm efficiently, similar to Agarwal *et al.* and Kaplan *et al.* [1, 10]. But unlike their
 76 algorithms which take $\Omega(n)$ time to find an augmenting path, we show that after $O(n \text{polylog } n)$
 77 preprocessing, an augmenting path can be found in $O(k \text{polylog } n)$ time. The key is to recycle
 78 (rather than rebuild) our data structures from one augmentation to the next. We refer to
 79 this idea as the *rewinding mechanism*.

80 Next in Sections 3, we obtain a $(1 + \varepsilon)$ -approximation algorithm for the geometric partial
 81 matching problem in \mathbb{R}^2 by providing an efficient implementation of the unit-capacity min-cost
 82 flow algorithm by Goldberg *et al.* [8].

83 **► Theorem 1.2.** *Given two point sets A and B in \mathbb{R}^2 each of size at most n , an integer*
 84 *$k \leq n$, and a parameter $\varepsilon > 0$, a $(1 + \varepsilon)$ -approximate min-cost matching of size k between A*
 85 *and B can be computed in $O((n + k\sqrt{k}) \text{polylog } n \cdot \log \varepsilon^{-1})$ time.*

86 The main challenge here is how to deal with the set of *dead nodes*, which neither have
 87 excess/deficit nor have flow passing through them, but still contribute to the size of the
 88 graph. We show that the number of *alive* nodes is only $O(k)$, and then represent the dead
 89 nodes implicitly so that the Hungarian search and computation of a blocking flow can be
 90 implemented in $O(k \text{polylog } n)$ time.

91 Finally in Section 4 we present a faster algorithm for the transportation problem in \mathbb{R}^2
 92 when the two point sets are unbalanced.

93 **► Theorem 1.3.** *Given two point sets A and B in \mathbb{R}^2 of sizes r and n respectively with*
 94 *$r \leq n$, along with supply-demand function $\phi : A \cup B \rightarrow \mathbb{Z}$, an optimal transportation map*
 95 *between A and B can be computed in $O(\min\{n^2, rn^{3/2}\} \text{polylog } n)$ time.*

96 Our result improves over the $O(n^2 \text{polylog } n)$ time algorithm in [2] for $r = o(\sqrt{n})$. As
 97 in [2], we also use the strongly polynomial uncapacitated minimum-cost flow algorithm by
 98 Orlin [13], but we additional ideas are needed to implement it faster for $r \leq \sqrt{n}$. Unlike in
 99 the case of matchings, the support of the transportation problem may have size $\Omega(n)$ even
 100 when r is a constant; so naïvely we can no longer spend time proportional to the size of
 101 the support of the transportation map. However, we ensure that the support is acyclic and
 102 use a data structure that can find an augmenting path in $O(r\sqrt{n} \text{polylog } n)$ time, assuming
 103 $r \leq \sqrt{n}$.

104 2 Minimum-Cost Partial Matchings using Hungarian Algorithm

105 In this section, we solve the geometric partial matching problem and prove Theorem 1.1 by
 106 implementing the Hungarian algorithm for partial matching in $O((n + k^2) \text{polylog } n)$ time.

A vertex v is *matched* by matching $M \subseteq E$ if v is the endpoint of some edge in M ; otherwise v is *unmatched*. Given a matching M , an *augmenting path* $\Pi = (a_1, b_1, \dots, a_\ell, b_\ell)$ is an odd-length path with unmatched endpoints (a_1 and b_ℓ) that alternates between edges outside and inside of M . The symmetric difference $M \oplus \Pi$ creates a new matching of size $|M| + 1$, called the *augmentation* of M by Π . The dual to the standard linear program for partial matching has dual variables for each vertex, called *potentials* π . Given potentials π , we can define the *reduced cost* on the edges to be $c_\pi(v, w) := c(v, w) - \pi(v) + \pi(w)$. Potentials π are *feasible* if the reduced costs are nonnegative for all edges in G . We say that an edge (v, w) is *admissible* under potentials π if $c_\pi(v, w) = 0$.

Fast implementation of Hungarian search. The Hungarian algorithm is initialized with $M = \emptyset$ and $\pi = 0$. Each iteration of the Hungarian algorithm augments M by an admissible augmenting path Π , discovered using a procedure called the *Hungarian search*. The algorithm terminates after k augmentations, when $|M| = k$; Ramshaw and Tarjan [14] showed that M is guaranteed to be an optimal partial matching.

The Hungarian search grows a set of *reachable vertices* X from all unmatched $v \in A$ using augmenting paths of admissible edges. Initially, X is the set of unmatched vertices in A . Let the *frontier* of X be the edges in $(A \cap X) \times (B \setminus X)$. X is grown by *relaxing* an edge (a, b) in the frontier: adding b into X , and also modifying potentials to make (a, b) admissible, preserve c_π on other edges within X , and keep π feasible on edges outside of X . Specifically, the algorithm relaxes the minimum-reduced-cost frontier edge (a, b) , and then raises $\pi(v)$ by $c_\pi(a, b)$ for all $v \in X$. If b is already matched, then we also relax the matching edge (a', b) and add a' into X . The search finishes when b is unmatched, and an admissible augmenting path now can be recovered.

In the geometric setting, we find the min-reduced-cost frontier edge using a dynamic *bichromatic closest pair* (BCP) data structure, as observed in [2, 18]. Given two point sets P and Q in the plane and a weight function $\omega : P \cup Q \rightarrow \mathbb{R}$, the BCP is two points $a \in P$ and $b \in Q$ minimizing the additively weighted distance $c(a, b) - \omega(a) + \omega(b)$. Thus, a minimum reduced-cost frontier edge is precisely the BCP of point sets $P = A \cap X$ and $Q = B \setminus X$, with $\omega = \pi$. Note that the “state” of this BCP is parameterized by X and π .

The dynamic BCP data structure by Kaplan *et al.* [10] supports point insertions and deletions in $O(\text{polylog } n)$ time and answers queries in $O(\log^2 n)$ time for our setting. Outside of potential updates, each relaxation in the Hungarian search requires one query, one deletion, and at most one insertion. As $|M| \leq k$ throughout, there are at most $2k$ relaxations in each Hungarian search, and the BCP can be used to implement each Hungarian search in $O(k \text{ polylog } n)$ time. Explained shortly, there is an existing technique to handle potential updates without performing BCP updates for each one.

Rewinding mechanism. We observe that exactly one vertex of A is newly matched after an augmentation. Thus (modulo potential changes), given the initial state of the BCP at the i -th Hungarian search, we can obtain the initial state for the $(i + 1)$ -th with a single BCP deletion operation.

If we remember the sequence of points added to X in the i -th Hungarian search, then at the start of the $(i + 1)$ -th Hungarian search we can *rewind* this sequence by applying the opposite insert/delete operation for each BCP update in reverse order to obtain the initial state of the i -th BCP. With one additional BCP delete, we have the initial state for the $(i + 1)$ -th BCP. The number of insertions/deletions is $O(k)$, bounded by the number of relaxations per Hungarian search, therefore we can recover, in $O(k \text{ polylog } n)$ time, the initial

BCP data structure for each Hungarian search beyond the first. We refer to this procedure as the *rewinding mechanism*.

«I tried to condense this; please check»

Potential updates. We modify a trick from Vaidya [18] to batch potential updates. Potentials are tracked with a *stored value* $\gamma(v)$, while the *true value* of $\pi(v)$ may have changed since $\gamma(v)$ was last recorded. This is done by aggregating potential changes into a variable δ , which is initially 0 at the very beginning of the algorithm. Whenever we would raise the potentials of all vertices in X , we raise δ by that amount instead. We maintain the following invariant: $\pi(v) = \gamma(v)$ for $v \notin X$, and $\pi(v) = \gamma(v) + \delta$ for $v \in X$.

At the beginning of the algorithm, X is empty and stored values are equal to true values. When $a \in A$ is added to X , we update its stored value to $\pi(a) - \delta$ for the current value of δ , and use that stored value as its BCP weight. Since the BCP weights are uniformly offset from $\pi(v)$ by δ , the pair reported by the BCP is still minimum. When $b \in B$ is added to X , we update its stored value to $\pi(b) - \delta$ (although it won't be added to a BCP set). When a vertex is removed from X (e.g. by augmentation or rewinding), we update the stored potential $\gamma(v) \leftarrow \pi(v) + \delta$, again for the current value of δ . Unlike [18], we do not reset δ across Hungarian searches.

There are $O(k)$ relaxations and thus $O(k)$ updates to δ per Hungarian search. $O(k)$ stored values are updated per rewinding, so the time spent on potential updates per Hungarian search is $O(k)$. Putting everything together, our implementation of the Hungarian algorithm runs in $O((n + k^2) \text{polylog } n)$ time. This proves Theorem 1.1.

3 Approximating Min-Cost Partial Matching through Cost-Scaling

In this section we describe an approximation algorithm for computing a min-cost partial matching. We reduce the problem to computing a min-cost circulation in a flow network (Section 3.1). We adapt the cost-scaling algorithm by Goldberg *et al.* [8] for computing min-cost flow of a unit-capacity network (Section 3.2). Finally, we show how their algorithm can be implemented in $O((n + k^{3/2}) \text{polylog}(n) \log(1/\epsilon))$ time in our setting (Section 3.3).

3.1 From matching to circulation

Given a bipartite graph G with vertex sets A and B , we construct a flow network $N = (V, \vec{E})$ in a standard way [14] so that a min-cost matching in G corresponds to a min-cost integral circulation in N .

Flow network. Each vertex in G becomes a node in N and each edge (a, b) in G becomes an arc $a \rightarrow b$ in N ; we refer to these nodes (resp. arcs) as *bipartite nodes* (resp. *bipartite arcs*). We also include a *source* node s and *sink* node t in N . For each $a \in A$, we add a *left dummy arc* $s \rightarrow a$ and for each $b \in B$ we add a *right dummy arc* $b \rightarrow t$. The cost $c(v \rightarrow w)$ of each arc $v \rightarrow w$ in N is equal to $c(v, w)$ if $v \rightarrow w$ is a bipartite arc and 0 if $v \rightarrow w$ is a dummy arc. All arcs in N have unit capacity.

Let $\phi : V \rightarrow \mathbb{Z}$ be an integral supply/demand function on nodes of N . The positive values of $\phi(v)$ are referred to as *supply*, and the negative values of $\phi(v)$ as *demand*. A *pseudoflow* $f : \vec{E} \rightarrow [0, 1]$ is a function on arcs of N . The *support* of f in N , denoted as $\text{supp}(f)$, is the set of arcs with positive flows: $\text{supp}(f) := \{v \rightarrow w \in \vec{E} \mid f(v \rightarrow w) > 0\}$. Given a pseudoflow f ,

194 the *imbalance* of a vertex (with respect to f) is

$$195 \quad \phi_f(v) := \phi(v) + \sum_{w \rightarrow v \in \vec{E}} f(w \rightarrow v) - \sum_{v \rightarrow w \in \vec{E}} f(v \rightarrow w).$$

196 We call positive imbalance *excess* and negative imbalance *deficit*; and vertices with positive
 197 (resp. negative) imbalance excess (resp. deficit) vertices. A vertex is *balanced* if it has zero
 198 imbalance. If all vertices are balanced, the pseudoflow is a *circulation*. The *cost* of a
 199 pseudoflow is defined to be

$$200 \quad c(f) := \sum_{v \rightarrow w \in \text{supp}(f)} c(v \rightarrow w) \cdot f(v \rightarrow w).$$

201 The *minimum-cost flow problem* (MCF) asks to find a circulation of minimum cost inside a
 202 given network.

203 If we set $\phi(s) = k$, $\phi(t) = k$, and $\phi(v) = 0$ for all $v \in A \cup B$, then an integral circulation
 204 f corresponds to a partial matching M of size k and vice versa. Moreover, $c(M) = c(f)$.
 205 Hence, the problem of computing a min-cost matching of size k in $G(A, B)$ transforms to
 206 computing an integral circulation in N . The following lemma will be useful for our algorithm.

207 ► **Lemma 3.1.** *Let N be the network constructed from $G(A, B)$ above.*

- 208 (i) *For any integral circulation g in N , the size of $\text{supp}(g)$ is at most $3k$.*
 209 (ii) *For any integral pseudoflow f in N with $O(k)$ excess, the size of $\text{supp}(f)$ is $O(k)$.*

210 3.2 A cost-scaling algorithm

211 Before describing the algorithm, we need to introduce a few more concepts.

212 **Residual network and admissibility.** If f is an integral pseudoflow (that is, $f(v \rightarrow w) \in \{0, 1\}$
 213 for every arc in \vec{E}), then each arc $v \rightarrow w$ in N is either *idle* with $f(v \rightarrow w) = 0$ or *saturated*
 214 with $f(v \rightarrow w) = 1$. *«define earlier?»*

215 Given a pseudoflow f , the *residual network* $N_f = (V, \vec{E}_f)$ is defined as follows. For each
 216 idle arc $v \rightarrow w$ in \vec{E} , we add a *forward* residual arc $v \rightarrow w$ in N_f . For each saturated arc $v \rightarrow w$
 217 in \vec{E} , we add a *backward* residual arc $w \rightarrow v$ in N_f . The set of residual arcs in N_f is therefore

$$218 \quad \vec{E}_f := \{v \rightarrow w \mid f(v \rightarrow w) = 0\} \cup \{w \rightarrow v \mid f(v \rightarrow w) = 1\}.$$

219 The cost of a forward residual arc $v \rightarrow w$ is $c(v \rightarrow w)$, while the cost of a backward residual arc
 220 is $w \rightarrow v$ is $-c(v \rightarrow w)$. Each arc in N_f also has unit capacity. By Lemma 3.1, N_f has $O(k)$
 221 backward arcs if f has $O(k)$ excess.

222 A *residual pseudoflow* g in N_f can be used to change f into a different pseudoflow on N
 223 by a process called *augmentation*. For simplicity, we only describe augmentation for the case
 224 where f, g are integer. Specifically, augmenting f by g produces a pseudoflow f' in N where

$$225 \quad f'(v \rightarrow w) = \begin{cases} 0 & w \rightarrow v \in \vec{E}_f \text{ and } g(w \rightarrow v) = 1 \\ 1 & v \rightarrow w \in \vec{E}_f \text{ and } g(v \rightarrow w) = 1 \\ f(v \rightarrow w) & \text{otherwise.} \end{cases}$$

226 Using LP duality for min-cost flow, we assign a *potential* $\pi(v)$ to each node v in N . The
 227 *reduced cost* of an arc $v \rightarrow w$ in N with respect to π is defined as

$$228 \quad c_\pi(v \rightarrow w) := c(v \rightarrow w) - \pi(v) + \pi(w).$$

Similarly we define the reduced cost of arcs in N_f : the reduced cost of a forward residual arc $v \rightarrow w$ in N_f is $c_\pi(v \rightarrow w)$, and the reduced cost of a backward residual arc $w \rightarrow v$ in N_f is $-c_\pi(v \rightarrow w)$. Abusing the notation, we also use c_π to denote the reduced cost of arcs in N_f .

«Unclear. Does $c_\pi(w \rightarrow v)$ have positive or negative value for a backward arc $w \rightarrow v$?»

The *dual feasibility constraint* asks that $c_\pi(v \rightarrow w) \geq 0$ holds for every arc $v \rightarrow w$ in \vec{E} ; potentials π that satisfy this constraint are said to be *feasible*. Suppose we relax the dual feasibility constraint to allow some small violation in the value of $c_\pi(v \rightarrow w)$. We say that a pair of pseudoflow f and potential π is *θ -optimal* [4, 17] if $c_\pi(v \rightarrow w) \geq -\theta$ for every residual arc $v \rightarrow w$ in \vec{E}_f . Pseudoflow f is *θ -optimal* if it is θ -optimal with respect to some potentials π ; potential π is *θ -optimal* if it is θ -optimal with respect to some pseudoflow f . Given a pseudoflow f and potentials π , a residual arc $v \rightarrow w$ in \vec{E}_f is *admissible* if $c_\pi(v \rightarrow w) \leq 0$. We say that a pseudoflow g in G_f is *admissible* if $g(v \rightarrow w) > 0$ only on admissible arcs $v \rightarrow w$, and $g(v \rightarrow w) = 0$ otherwise.² We will use the following well-known property of θ -optimality.

► **Lemma 3.2.** *Let f be an θ -optimal pseudoflow in N and let g be an admissible pseudoflow in N_f . Then f augmented by g is also θ -optimal in N .*

Using Lemma 3.1, the following lemma can be proved about θ -optimality:

► **Lemma 3.3.** *Let f be a θ -optimal integer circulation in N , and f^* be an optimal integer circulation for N . Then, $c(f) \leq c(f^*) + 6k\theta$.*

Estimating the value of $c(f^*)$. We now describe a procedure for estimating $c(f^*)$ within a polynomial factor, which will be useful in setting the scaling parameters of the cost-scaling algorithm.

Let T be a minimum spanning tree of $A \cup B$ under the cost function $c(\cdot)$. Let e_1, e_2, \dots, e_{n-1} be the edges of T sorted in nondecreasing order of length; in other words, $c(e_i) \leq c(e_{i+1})$. Let T_i be the forest consisting of the vertices of $A \cup B$ and e_1, \dots, e_i . We call a matching M of $G(A, B)$ *intra-cluster* if both endpoints of every edge in M lie in the same connected component of T_i . We define i^* to be the smallest index i such that there exists an intra-cluster matching of size k in T_{i^*} . Set $\bar{\theta} := n^q \cdot c(e_{i^*})$. The following lemma will be used by our cost-scaling algorithm:

► **Lemma 3.4.** (i) *The value of i^* can be computed in $O(n \log n)$ time.*
(ii) *$c(e_{i^*}) \leq c(f^*) \leq \bar{\theta}$.*
(iii) *There is a $\bar{\theta}$ -optimal circulation in the network N with respect to the 0 potential $\pi = 0$, assuming $\phi(s) = k$, $\phi(t) = -k$, and $\phi(v) = 0$ for all $v \in A \cup B$.*

Set $\underline{\theta} := \frac{\varepsilon}{6k} \cdot c(e_{i^*})$. As a consequence of Lemmas 3.4ii and 3.3, we have:

► **Corollary 3.5.** *The cost of a $\underline{\theta}$ -optimal integral circulation in N is at most $(1 + \varepsilon)c(f^*)$.*

We are now ready to describe our algorithm.

Overview of the algorithm. We closely follow the algorithm of Goldberg *et al.* [8]. The algorithm works in rounds. In the beginning of each round, we fix a *cost scaling parameter* θ and maintain potentials π with the following property:

«would like to use a better/neater/cleaner environment to represent invariant»

² The same admissibility/feasibility definitions will be used later in Section 4. However, the algorithm in Section 4 maintains a 0-optimal f and therefore admissible residual arcs always have $c_\pi(v \rightarrow w) = 0$.

(*) There exists a 2θ -optimal integral circulation in N with respect to π .

For the initial round, we set $\theta \leftarrow \bar{\theta}$ and $\pi \leftarrow 0$. By Lemma 3.4(iii), (*) is satisfied initially. Each round of the algorithm consists of two stages. In the first stage, called *scale initialization* (SCALE-INIT) computes a θ -optimal pseudoflow f . The second stage, called *refinement* (REFINE) converts f into a θ -optimal (integral) circulation g . In both stages, π is updated as necessary. If $\theta \leq \underline{\theta}$, we return g . Otherwise, we set $\theta \leftarrow \theta/2$ and start the next round. Note that (*) is satisfied in the beginning of each round.

By Corollary 3.5, when the algorithm terminates, it returns an integral circulation \tilde{f} in N of cost at most $(1 + \varepsilon)c(f^*)$, which corresponds to a $(1 + \varepsilon)$ -approximate min-cost matching of size k in G . The algorithm terminates in $\log_2(\bar{\theta}/\underline{\theta}) = O(\log(n/\varepsilon))$ rounds.

Next, we describe the two stages in detail.

Scale initialization. In the first round, we compute a $\bar{\theta}$ -optimal pseudoflow by simply setting $f(v \rightarrow w) = 0$ for all arcs in \bar{E} . For subsequent rounds, we adjust the potential and flow in N as follows: we raise the potential of all nodes in A by θ , those in B by 2θ , and of t by 3θ . The potential of s remains unchanged. Since the reduced cost of every forward arc in N_f after the previous round is at least -2θ , the above step increases the reduced cost of all forward arcs by θ , and the reduced cost of all forward arcs is at least $-\theta$.

Next, for each backward arc $w \rightarrow v$ in N_f with $c_\pi(w \rightarrow v) < -\theta$, we set $f(v \rightarrow w) = 0$ (that is, make arc $v \rightarrow w$ idle), which replaces the backward arc $w \rightarrow v$ in N_f with a forward arc of positive reduced cost. After this step, the resulting pseudoflow must be θ -optimal as all arcs of N_f have reduced cost at least $-\theta$.

The desaturation of each backward arc creates one unit of excess. Since there are at most $3k$ backward arcs, the total excess in the resulting pseudoflow is at most $3k$. There are $O(n)$ potential updates and $O(k)$ arcs on which flow might change, therefore the time needed for SCALE-INIT is $O(n)$.

Refinement. The procedure REFINE converts a θ -optimal pseudoflow with $O(k)$ excess into a θ -optimal circulation, using a primal-dual augmentation algorithm. A path in N_f is an *augmenting path* if it begins at an excess vertex and ends at a deficit vertex. We call an admissible pseudoflow g in the residual network N_f an *admissible blocking flow* if g saturates at least one arc in every admissible augmenting path in N_g . In other words, there is no admissible excess-deficit path in the residual network after augmentation by g . Each iteration of REFINE finds an admissible blocking flow to be added to the current pseudoflow in two steps:

1. *Hungarian search*: a Dijkstra-like search that begins at the set of excess vertices and raises potentials until there is an excess-deficit path of admissible arcs in N_f .
2. *Augmentation*: using depth-first search through the set of admissible arcs of N_f , construct an admissible blocking flow. It suffices to repeatedly extract admissible augmenting paths until no more admissible excess-deficit paths remain.

The algorithm repeats these steps until the total excess becomes zero. The following lemma bounds the number of iterations in the REFINE procedure at each scale.

► **Lemma 3.6.** *Let θ be the scaling parameter and π_0 the potential function at the beginning of a round, such that there exists an integral 2θ -optimal circulation with respect to π_0 . Let f be a θ -optimal pseudoflow with excess $O(k)$. Then REFINE terminates within $O(\sqrt{k})$ iterations.*

Proof. We sketch the proof, which is adapted from [8]. Let f_0 be the assumed 2θ -optimal integral circulation with respect to π_0 , and let π be the potentials maintained and changing

during REFINE. Let $d(v) = (\pi(v) - \pi_0(v))/\theta$, i.e. the increase in potential at v in units of θ . We divide the iterations of REFINE into two phases: before and after every (remaining) excess vertex has $d(v) \geq \sqrt{k}$. Each Hungarian search raises excess potentials by at least θ , since we use blocking flows. Thus, the first phase lasts at most \sqrt{k} iterations.

At the start of the second phase, consider the set of arcs $E^+ = \{v \rightarrow w \in \vec{E} \mid f(v \rightarrow w) < f_0(v \rightarrow w)\}$. One can argue that the remaining excess with respect to f is bounded above by the size of any cut separating the excess and deficit vertices. The proof examines cuts $Y_i = \{v \mid d(v) > i\}$ for $0 \leq i \leq \sqrt{k}$. By θ -optimality of f and 2θ -optimality of f_0 , one can show that each arc in E^+ crosses at most 3 cuts. Furthermore, the size of E^+ is $O(k)$, bounded by the support size of f and f_0 . Averaging, there is a cut among the Y_i of size at most $3k/\sqrt{k}$, so the total excess remaining is $O(\sqrt{k})$. Each iteration of REFINE eliminates at least one unit of excess, so the number of second phase iterations is also at most $O(\sqrt{k})$. ◀

In the next subsection we show that after $O(n \text{ polylog } n)$ preprocessing, an iteration of REFINE can be performed in $O(k \text{ polylog } n)$ time (cf. Lemma 3.8). By Lemma 3.6 and the fact the algorithm terminates in $O(\log(n/\varepsilon))$ rounds, the overall running time of the algorithm is $O((n + k^{3/2}) \text{ polylog } n \log(1/\varepsilon))$, as claimed in Theorem 1.2.

3.3 Fast implementation of refinement

We now describe a fast implementation of the refinement stage. The Hungarian search and augmentation steps are similar: each traversing through the residual network using admissible arcs starting from the excess vertices. Due to lack of space, we only describe the Hungarian search process.

At a high level, let X be the subset of nodes visited by the Hungarian search so far. Initially X is the set of excess nodes. At each step, the algorithm finds a minimum reduced cost arc $v \rightarrow w$ in N_f from X to $V \setminus X$. If $v \rightarrow w$ is not admissible, the potential of all nodes in X is increased by $\lceil c_\pi(v \rightarrow w)/\theta \rceil$ to make $v \rightarrow w$ admissible. If w is a deficit node, the search terminates. Otherwise, w is added to X and the search continues.

Implementing the Hungarian search efficiently is more difficult than in Section 2 because (a) excess nodes may show up in A as well as in B , (b) a balanced node may become imbalanced later in the rounds, and (c) the potential of excess nodes may be non-uniform. We therefore need a more complex data structure.

We call a node v of N *dead* if $\phi_f(v) = 0$ and no arc of $\text{supp}(f)$ is incident to it; otherwise v is *alive*. We note that s and t are always alive. Let A^* (resp. B^* , X^*) denote the set of alive nodes of A (resp. B , X). There are only $O(k)$ alive nodes, as each can be charged to its adjoining $\text{supp}(f)$ arcs or its excess/deficit. We treat alive and dead nodes separately to implement the Hungarian search efficiently. By definition, dead vertices only adjoin forward arcs in N_f . Thus, the in-degree (resp. out-degree) of a node in $(A \setminus A^*)$ (resp. $(B \setminus B^*)$) is 1, and any path passing through a dead vertex has a subpath of the form $s \rightarrow v \rightarrow b$ for some $b \in B$ or $a \rightarrow v \rightarrow t$ for some $a \in A$. Consequently, a path in N_f may have at most two consecutive dead nodes, and in the case of two consecutive dead nodes there is a subpath of the form $s \rightarrow v \rightarrow w \rightarrow t$ where $v \in (A \setminus A^*)$ and $w \in (B \setminus B^*)$. We will call such paths, from an alive node to an alive node with only dead interior nodes, *alive-alive paths*. We extend the notions of reduced cost and admissibility to alive-alive paths, where the reduced cost of the path is the sum of its edges. Since reduced costs telescope, the reduced cost of an alive-alive path depends only on the potentials at its (alive) endpoints.

Using this observation, we implement the Hungarian search to “skip over” dead nodes, while logically exploring the same alive nodes in the same order. Alive-alive paths may have

length 1 (no dead interior nodes), 2, or 3. At each step, we find a minimum reduced cost alive-alive path Π from an alive node of X to an alive node of $V \setminus X$, and add the nodes of this path into X in a single step. We update potentials in X according to the reduced cost of the path. There are $O(k)$ alive nodes, so the number of minimization queries per Hungarian search is $O(k)$.

We find the minimum reduced cost alive-alive path of length 1, 2, and 3, and then choose the cheapest among them. We now describe a data structure for each path length. For each data structure, our “time budget” per Hungarian search is $O(k \text{ polylog } n)$.

Finding length-1 paths. This data structure finds a minimum reduced cost arc from an alive node of X to an alive node of $V \setminus X$. There are $O(k)$ backward arcs, so the minimum among backward arcs can be maintained explicitly in a priority queue on c_π and retrieved in $O(1)$ time.

There are three types of forward arcs: $s \rightarrow a$ for some $a \in A^*$, $b \rightarrow t$ for some $b \in B^*$, and bipartite arc $a \rightarrow b$ with two alive endpoints. Edges of the first (resp. second) type can be found by maintaining $A^* \setminus X$ (resp. $B^* \cap X$) in a priority queues on π , but should only be queried if $s \in X$ (resp. $t \notin X$). The cheapest arc of the third type is an (additively weighted) bichromatic closest pair (BCP) between $A^* \cap X$ and $B^* \setminus X$, with reduced cost as the pair distance $\langle\langle \text{potentials?} \rangle\rangle$. We thus maintain $A^* \cap X$, $B^* \setminus X$ in a dynamic BCP data structure [10] on which insertions/deletions can be performed in $O(\text{polylog } k)$ time.

Finding length-2 paths. We describe how to find the cheapest path of the form $s \rightarrow v \rightarrow b$ where v is dead and $b \in B^*$. A cheapest path of the form $a \rightarrow v \rightarrow t$ can be found similarly. As for length-1 paths, we only query paths originating from s if $s \in X$, and only query paths ending at t if $t \notin X$.

Note that $c_\pi(s \rightarrow v \rightarrow b) = c(v, b) + \pi(b) - \pi(s)$. Since $\pi(s)$ is common in all such paths, it suffices to find the pair minimizing

$$\min_{v \in (A \setminus A^*), w \in B^* \setminus X} c(v, w) + \pi(w)$$

This is done by maintaining a dynamic BCP data structure between $(A \setminus A^*)$ and $B^* \setminus X$ with the cost of a pair (v, w) being $c(v, w) + \pi(w)$. We may require an update operation for each alive node added to X during the Hungarian search, of which there are $O(k)$, so the time spent during a search is $O(k \text{ polylog } n)$.

Since the size of $(A \setminus A^*)$ is at least $r - k$, we cannot construct this BCP from scratch at the beginning of each iteration of Hungarian search. To resolve this, we use the idea of rewinding from Section 2, with a slight twist. There are now *two* ways that the initial BCP may change across consecutive Hungarian searches: (1) the initial set X may change as vertices lose excess through augmentation, and (2) the set of dead A vertices may change; that is, when flow is augmented across a vertex of $(A \setminus A^*)$. The first is identical to the situation in Section 2; the number of excess depletions is $O(k)$ over the course of REFINES. For the second, the alive/dead status of a node can change only if the blocking flow found in the augmentation process passes through it. By Lemma 3.7(2) below, there are $O(k)$ such changes. Thus, updating $(A \setminus A^*)$ for the BCP (after augmentation) can be done in $O(k \text{ polylog } n)$ time for each Hungarian search.

Finding length-3 paths. We now describe how to find the cheapest path of the form $s \rightarrow v \rightarrow w \rightarrow t$ where $v \in (A \setminus A^*)$ and $w \in (B \setminus B^*)$. Note that $c_\pi(s \rightarrow v \rightarrow w \rightarrow t) = c(v \rightarrow w) -$

404 $\pi(s) + \pi(t)$. Hence, we need to find a BCP between $(A \setminus A^*)$ and $(B \setminus B^*)$, where the cost
 405 of a pair (v, w) is simply $c(v, w)$. This can be done by maintaining a dynamic BCP data
 406 structure similar to the case of length-2 paths.

407 The BCP sets have no dependence on X — the only updates required come from
 408 membership changes to A^* or B^* , after an augmentation. Applying Lemma 3.7(2) again,
 409 there are $O(k)$ alive/dead updates caused by an augmentation, so the time for these updates
 410 per Hungarian search is $O(k \text{ polylog } n)$.

411 **Updating potentials.** The Hungarian search periodically raises the potentials for all nodes
 412 in X , and we need to implement this efficiently for the data structures above. Note that the
 413 data structures above do not utilize potentials of dead nodes. Potentials for alive nodes can
 414 be updated in a batched fashion using the method in Section 2.

415 For dead nodes, we ignore their potentials entirely and instead recover “valid” potentials
 416 for them once they switch from dead to alive (and additionally, at the end of a round).
 417 Specifically, for a newly alive $a \in A$ we set $c_\pi(a) \leftarrow c_\pi(s)$ and for newly-alive $b \in B$ we set
 418 $c_\pi(b) \leftarrow c_\pi(t)$. Formally, we will say an alive-alive path is *strongly admissible* under π if
 419 all of its arcs are admissible under π . For correctness, we need to show that our choice of
 420 recovered potentials (1) preserves θ -optimality and (2) makes any admissible alive-alive path
 421 from before augmentation strongly admissible. It is a straightforward calculation to verify
 422 that both properties hold.

423 The following lemmas are crucial to analyzing the running time of the Hungarian search.

- 424 ► **Lemma 3.7.** 1. *Both the Hungarian search and augmentation step explore $O(k)$ nodes.*
 425 2. *The blocking flow found by augmentation is incident to $O(k)$ vertices.*

426 Augmentation can also be implemented in $O(k \text{ polylog } n)$ time, after $O(n \text{ polylog } n)$
 427 preprocessing, using similar data structures. We thus obtain the following:

- 428 ► **Lemma 3.8.** *After $O(n \text{ polylog } n)$ preprocessing, each iteration of REFINES can be imple-*
 429 *mented in $O(k \text{ polylog } n)$ time.*

430 **4 Transportation Algorithm**

431 Given two point sets A and B in \mathbb{R}^2 of sizes r and n respectively (with $r \leq \sqrt{n}$) and
 432 a supply-demand function $\phi : A \cup B \rightarrow \mathbb{Z}$ as defined in the introduction, we present an
 433 $O(rn^{3/2} \text{ polylog } n)$ time algorithm for computing an optimal transport map between A and
 434 B . By applying this algorithm in the case of $r \leq \sqrt{n}$ and the one in [2] when $r > \sqrt{n}$, we
 435 prove Theorem 1.3. We use a standard reduction to the uncapacitated min-cost flow problem
 436 (see e.g. [13]) and use Orlin’s algorithm [13] as well as some of the ideas from [2] **«The full**
 437 **version of the transportation paper is still not available anywhere.»** for implementing it
 438 faster in geometric settings. We first present an overview of the algorithm and then describe
 439 its fast implementation to achieve the desired running time.

440 **4.1 Overview of the algorithm**

442 Orlin’s algorithm follows an excess-scaling paradigm and the primal-dual framework. It
 443 maintains a *scale parameter* Δ , a flow function f , and potentials π on the nodes. Initially
 444 $\Delta = \phi(A)$, $f = 0$, and $\pi = 0$. We fix a constant parameter $\alpha \in (0.5, 1)$. A node v is called
 445 *active* if $|\phi_f(v)| \geq \alpha\Delta$. At each step, using the Hungarian search, the algorithm finds an
 446 admissible active-excess-to-active-deficit path in the residual network and pushes a flow of Δ

447 along this path.³ It repeats this step as long as there are both active excess and active deficit
 448 nodes. When one of these sets becomes empty, Δ is halved. The sequence of augmentations
 449 with a fixed value of Δ is called an *excess scale*.

451 The algorithm also performs two preprocessing steps at the beginning of each excess scale.
 452 If $f(v \rightarrow w) \geq 3n\Delta$, $v \rightarrow w$ is contracted to a single node z with $\phi(z) = \phi(v) + \phi(w)$.⁴ If there
 453 are no active excess nodes and $f(v \rightarrow w) = 0$ for all arcs, Δ is lowered to $\max_v \phi(v)$.

454 When the algorithm terminates, it has found an optimal circulation in the contracted
 455 network. We use the algorithm described in [2] to recover an optimal circulation in the
 456 original network. Orlin showed that the algorithm terminates within $O(n \log n)$ scales
 457 and performs a total of $O(n \log n)$ augmentations. In the next subsection, we describe an
 458 algorithm that after $O(n \text{ polylog } n)$ preprocessing can perform a Hungarian search (i.e. find
 459 an excess-deficit admissible path) in $O(r\sqrt{n} \text{ polylog } n)$ amortized time. Summing this cost
 460 over all augmentations, we obtain the desired running time.

461 4.2 An efficient implementation

462 Recall in the previous sections that we could bound the running time of the Hungarian
 463 search by the size of $\text{supp}(f)$. Here, the number of active excess/deficit nodes at any scale
 464 is $O(r)$, and the length of an augmenting path is also $O(r)$. Therefore one might hope
 465 to find an augmenting path in $O(r \text{ polylog } n)$ time, by adapting the algorithms described
 466 in Sections 2 and 3. The challenge to this approach is that $\text{supp}(f)$ may have $\Omega(n)$ size,
 467 therefore an algorithm which runs in $O(|\text{supp}(f)|)$ time is no longer sufficient. Still, we
 468 manage to implement a Hungarian search in time roughly $r\sqrt{n}$, by exploiting a few properties
 469 of $\text{supp}(f)$ as described below.

470 We note that each arc of $\text{supp}(f)$ is admissible, i.e., its reduced cost is 0, so we add an
 471 arc of $\text{supp}(f)$ as soon as possible when it arrives in $X \times (V \setminus X)$. This strategy ensures the
 472 following crucial property.

473 ► **Lemma 4.1.** *If the arcs of $\text{supp}(f)$ are added as soon as possible, $\text{supp}(f)$ is acyclic.*

474 Next, similar to Section 3, we call a node $u \in V$ *alive* if u is an active excess/deficit node
 475 or if u is incident on an arc of $\text{supp}(f)$. Otherwise, u is called *dead*. Unlike Section 3, once a
 476 node becomes alive it cannot become dead. Furthermore, a dead node may become alive
 477 only in the beginning of a scale (after we have reduced the value of Δ). Unlike Section 3, an
 478 augmenting path cannot pass through a dead node. Therefore, we can ignore all dead nodes
 479 during Hungarian search, work with only alive nodes, and update the set of alive nodes in
 480 the beginning of a scale. For a subset $S \subseteq V$ of nodes, we use S^* to denote the set of alive
 481 nodes in S .

482 Let $B^* \subseteq B^*$ be the set of nodes that are either active excess/deficit nodes or that are
 483 incident on at least two arcs of $\text{supp}(f)$. Lemma 4.1 implies the following:

484 ► **Lemma 4.2.** $|B^*| = O(r)$.

485 We can therefore find the min reduced-cost arc $(X \cap A^*) \times (B^* \setminus X)$ using a BCP data
 486 structure as in Section 2, along with lazy potential updates and the rewinding mechanism.
 487 The total time spent by Hungarian search on the nodes of B^* will be $O(r \text{ polylog } n)$. We
 488 subsequently focus on handling $B^* \setminus B^*$.

441 ³ Note that this augmentation may convert an excess node into a deficit node.

450 ⁴ Intuitively, $f(v \rightarrow w)$ is so high that future scales cannot deplete the flow on $v \rightarrow w$.

Handling $B^\star \setminus B^\star$. We now describe how we query a min reduced-cost arc between $(X \cap A^\star)$ and $B^\star \setminus (B^\star \cup X)$. Each node $b \in B^\star \setminus B^\star$ is incident on exactly one arc of $\text{supp}(f)$ (i.e., there is only one outgoing arc from b). We partition these nodes into clusters depending on their neighbor in N_f . That is, for a node $a \in A^\star$, let $B_a = \{b \in B^\star \setminus B^\star \mid a \rightarrow b \in \text{supp}(f)\}$. We refer to B_a as the *star* of a .

The crucial observation is that a is the only node in N_f reachable from each $b \in B_a$, so once the Hungarian search reaches a node of B_a and thus a (recall we prioritize adding $\text{supp}(f)$ arcs), the Hungarian search need not visit any other nodes of B_a , as they will only lead to a . Hence, as soon as one node of B_a is reached, all other nodes of B_a are discarded from further consideration. Using this observation, we handle $B^\star \setminus B^\star$ as follows.

We classify each $a \in A^\star$ as *light* or *heavy*. If a is classified as heavy then $|B_a| \geq \sqrt{n}$, and if a is classified as light then $|B_a| \leq 2\sqrt{n}$. Note that if $|B_a| \in [\sqrt{n}, 2\sqrt{n}]$, then a may be classified as light or heavy. We allow this flexibility to allow re-classification in a lazy manner. Namely, a light node is re-classified heavy once $|B_a| > 2\sqrt{n}$, and a heavy node is re-classified light once $|B_a| < \sqrt{n}$. This scheme ensures that the star of a has gone through at least \sqrt{n} updates (insertion/deletion of nodes) between two successive re-classifications, and these updates will pay for the time spent in updating the data structure when a is re-classified — see below.

For each heavy node $a \in A^\star \setminus X$, we maintain a BCP data structure between B_a and $X \cap A^\star$. Next, for all light $a \in A^\star \setminus X$, we collect their stars into a single set $B^< = \bigcup_{a \text{ light}} B_a$. We maintain a BCP data structure between $B^<$ and $A^\star \cap X$. Thus, we maintain at most r different BCP data structures for stars.

Using these data structures, we can compute a min reduced-cost arc between $A^\star \cap X$ and $B^\star \setminus (B^\star \cup X)$. Once an arc $v \rightarrow w \in (A^\star \cap X) \times (B^\star \setminus (B^\star \cup X))$ is added such that $w \in B_a$ for some star B_a , then we also add a to X . If a is light, then we delete B_a from $B^<$ and update the BCP data structure for $B^<$. If a is heavy, then we stop querying the BCP data structure on B_a for the remainder of the search. Finally, since a becomes a part of X , a is added to all $O(r)$ BCP data structures.

Recall that $r \leq \sqrt{n}$ by assumption. Adding the arc $v \rightarrow w$ thus involves performing $O(\sqrt{n})$ insertion/deletion operations in various BCP data structures, thereby taking $O(\sqrt{n} \text{ polylog } n)$ time.

Putting it together. While proof is omitted, the following lemma bounds the running time of the Hungarian search.

► **Lemma 4.3.** *Assuming all BCP data structures are initialized correctly, the Hungarian search terminates within $O(r)$ steps, and takes $O(r\sqrt{n} \text{ polylog } n)$ time.*

Once an augmenting path is found and the augmentation is performed, the set of excess/deficit nodes and $\text{supp}(f)$ arcs change. We thus need to update the set B^\star , B_a , and $B^<$. This can be accomplished in $O(r \text{ polylog } n)$ amortized time. When we begin a new Hungarian search, we use the rewinding mechanism to set various BCP data structures in the right initial state. Finally, when we move from one scale to another, we also update the sets A^\star and B^\star . Omitting all the details, we conclude the following.

► **Lemma 4.4.** *Each Hungarian search can be performed in $O(r\sqrt{n} \text{ polylog } n)$ time.*

Since there are $O(n \log n)$ augmentations and the flow in the original network can be recovered from that in the contracted network in $O(n \log n)$ time [2], the total running time of the algorithm is $O(rn^{3/2} \text{ polylog } n)$, as claimed in Theorem 1.3.

534 ——— References ———

- 535 1 Pankaj K. Agarwal, Alon Efrat, and Micha Sharir. Vertical decomposition of shallow levels
536 in 3-dimensional arrangements and its applications. *SIAM J. Comput.* 29(3):912–953, 1999.
537 <https://doi.org/10.1137/S0097539795295936>.
- 538 2 Pankaj K. Agarwal, Kyle Fox, Debmalya Panigrahi, Kasturi R. Varadarajan, and Allen
539 Xiao. Faster algorithms for the geometric transportation problem. *Proc. 33rd Int. Sympos.*
540 *Comput. Geom. (SoCG)*, 7:1–7:16, 2017. [https://doi.org/10.4230/LIPIcs.SoCG.2017.](https://doi.org/10.4230/LIPIcs.SoCG.2017.7)
541 7).
- 542 3 Pankaj K. Agarwal and Kasturi R. Varadarajan. A near-linear constant-factor approxima-
543 tion for Euclidean bipartite matching? *Proc. 20th Annu. Sympos. Comput. Geom. (SoCG)*,
544 247–252, 2004. <https://doi.org/10.1145/997817.997856>.
- 545 4 D. Bertsekas and D. El Baz. Distributed asynchronous relaxation methods for convex
546 network flow problems. *SIAM J. Control and Opt.* 25(1):74–85, 1987. [https://doi.org/](https://doi.org/10.1137/0325006)
547 10.1137/0325006).
- 548 5 Ran Duan, Seth Pettie, and Hsin-Hao Su. Scaling algorithms for weighted matching in gen-
549 eral graphs. *ACM Trans. Algorithms* 14(1):8:1–8:35, 2018. [https://doi.org/10.1145/](https://doi.org/10.1145/3155301)
550 3155301).
- 551 6 Shimon Even and Robert E. Tarjan. Network flow and testing graph connectivity. *SIAM*
552 *J. Comput.* 4(4):507–518, 1975. <https://doi.org/10.1137/0204043>.
- 553 7 Harold N. Gabow and Robert E. Tarjan. Faster scaling algorithms for network problems.
554 *SIAM J. Comput.* 18(5):1013–1036, 1989. <https://doi.org/10.1137/0218069>.
- 555 8 Andrew V. Goldberg, Sagi Hed, Haim Kaplan, and Robert E. Tarjan. Minimum-cost
556 flows in unit-capacity networks. *Theoret. Comput. Sci.* 61(4):987–1010, 2017. <https://doi.org/10.1007/s00224-017-9776-7>.
- 558 9 John E. Hopcroft and Richard M. Karp. An $n^{5/2}$ algorithm for maximum matchings in
559 bipartite graphs. *SIAM J. Comput.* 2(4):225–231, 1973. [https://doi.org/10.1137/](https://doi.org/10.1137/0202019)
560 0202019).
- 561 10 Haim Kaplan, Wolfgang Mulzer, Liam Roditty, Paul Seiferth, and Micha Sharir. Dy-
562 namic planar Voronoi diagrams for general distance functions and their algorithmic applic-
563 ations. *Proc. 28th Annu. ACM-SIAM Sympos. Discrete Algorithms (SODA)*, 2495–2504,
564 2017. <https://doi.org/10.1137/1.9781611974782.165>.
- 565 11 Harold W. Kuhn. The Hungarian method for the assignment problem. *Naval Research*
566 *Logistics (NRL)* 2(1-2):83–97. John Wiley & Sons, 1955.
- 567 12 Aleksander Mądry. Navigating central path with electrical flows: From flows to matchings,
568 and back. *54th Annu. IEEE Sympos. Found. Comput. Sci. (FOCS)*, 253–262, 2013. <https://doi.org/10.1109/FOCS.2013.35>.
- 570 13 James B. Orlin. A faster strongly polynomial minimum cost flow algorithm. *Operations*
571 *Research* 41(2):338–350, 1993. <https://doi.org/10.1287/opre.41.2.338>.
- 572 14 Lyle Ramshaw and Robert E. Tarjan. A weight-scaling algorithm for min-cost imper-
573 fect matchings in bipartite graphs. *Proc. 53rd Annu. IEEE Sympos. Found. Comput. Sci.*
574 *(FOCS)*, 581–590, 2012. <https://doi.org/10.1109/FOCS.2012.9>.
- 575 15 R. Sharathkumar and Pankaj K. Agarwal. Algorithms for the transportation problem in
576 geometric settings. *Proc. 23rd Annu. ACM-SIAM Sympos. Discrete Algorithms (SODA)*,
577 306–317, 2012. <https://dl.acm.org/citation.cfm?id=2095116.2095145>.
- 578 16 R. Sharathkumar and Pankaj K. Agarwal. A near-linear time ϵ -approximation algorithm for
579 geometric bipartite matching. *Proc. 44th Annu. ACM Sympos. Theory Comput. (STOC)*,
580 385–394, 2012. <https://doi.org/10.1145/2213977.2214014>.
- 581 17 Éva Tardos. A strongly polynomial minimum cost circulation algorithm. *Combinatorica*
582 5(3):247–256, 1985. <https://doi.org/10.1007/BF02579369>.

- 583 18 Pravin M. Vaidya. Geometry helps in matching. *SIAM J. Comput.* 18(6):1201–1225, 1989.
584 [⟨https://doi.org/10.1137/0218080⟩](https://doi.org/10.1137/0218080).
- 585 19 Kasturi R. Varadarajan. A divide-and-conquer algorithm for min-cost perfect matching in
586 the plane. *Proc. 39th Annu. IEEE Sympos. Found. Comput. Sci. (FOCS)*, 320–331, 1998.
587 [⟨https://doi.org/10.1109/SFCS.1998.743466⟩](https://doi.org/10.1109/SFCS.1998.743466).